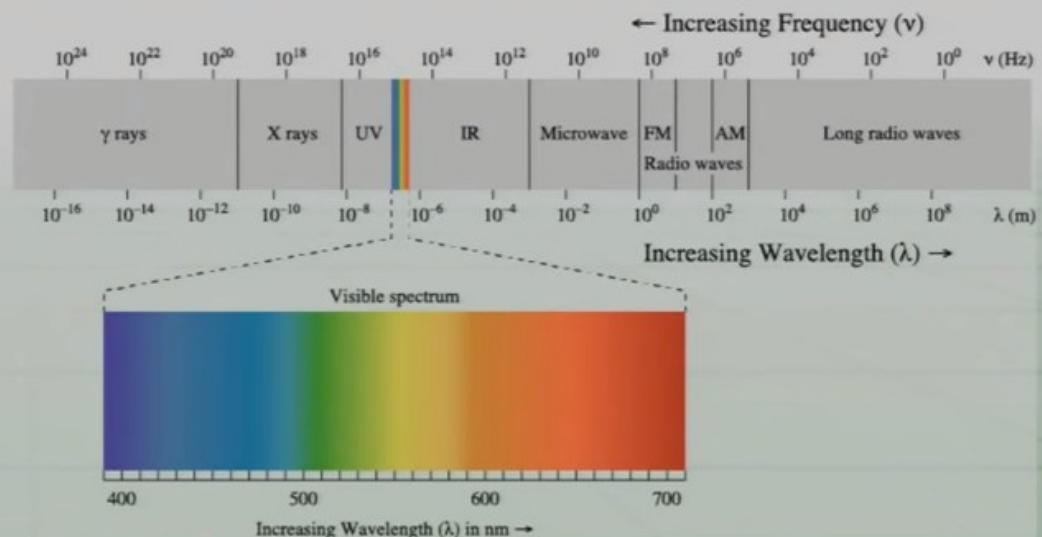
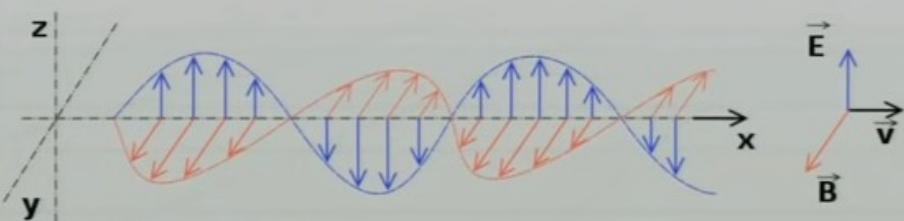
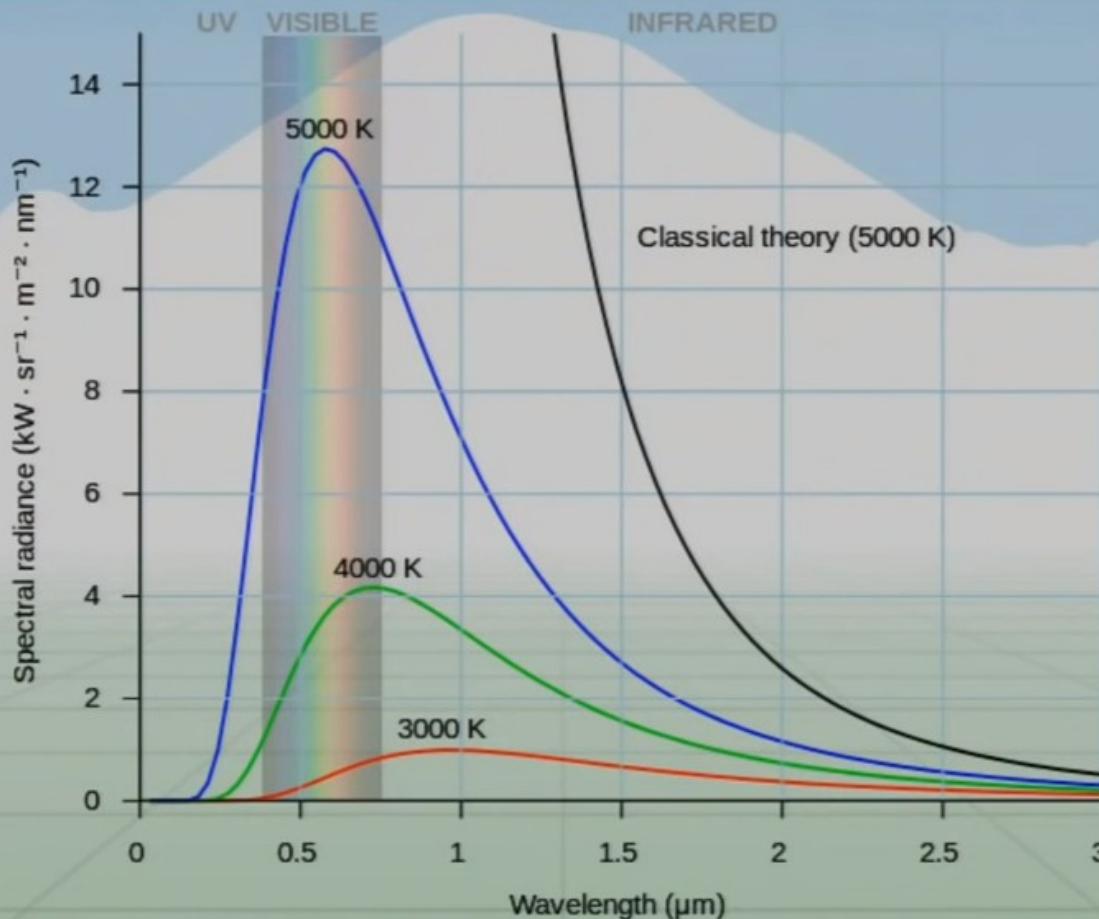


So what are we looking at anyway?

- Electromagnetic radiation
 - Wave? Particle?
- Photon - single particle of light
- Visible light: ~400-700 nanometers
- Why that range?

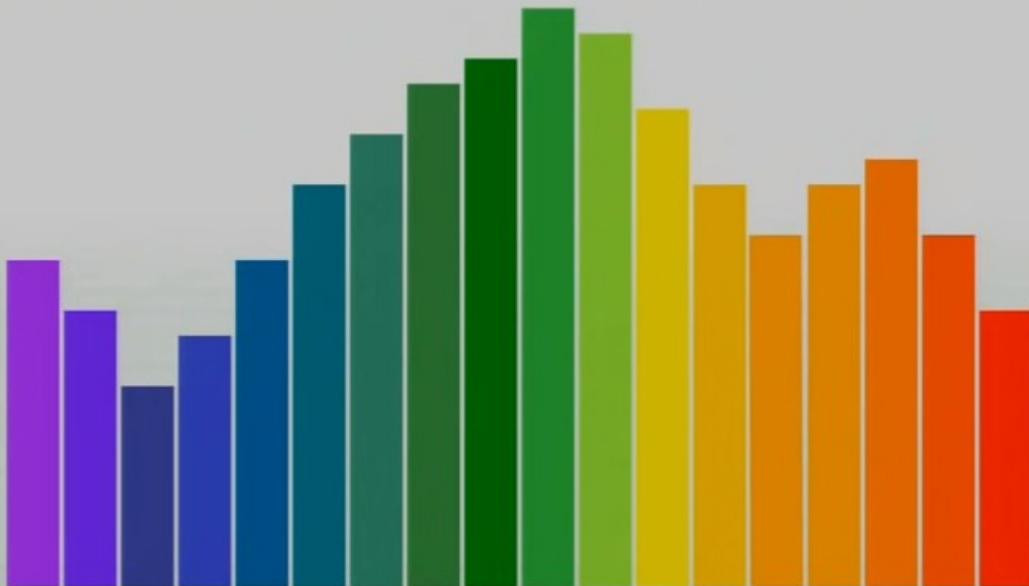


Visible light and the sun

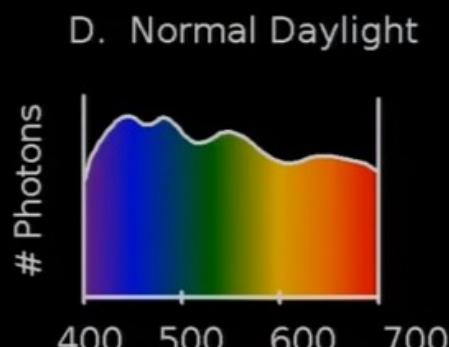
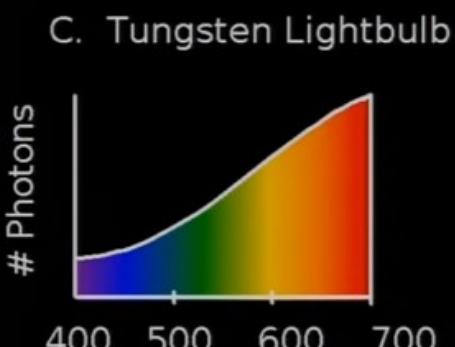
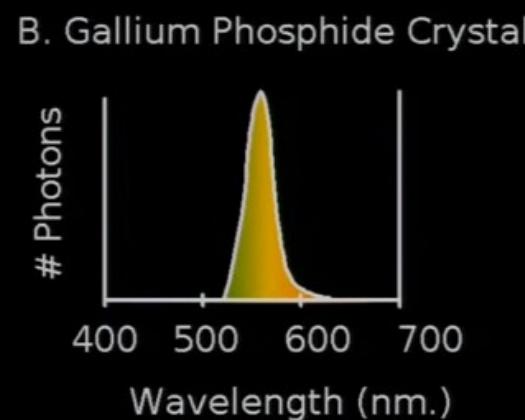
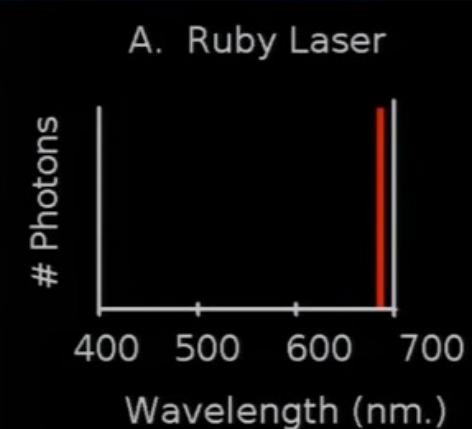


Light is a combination of waves

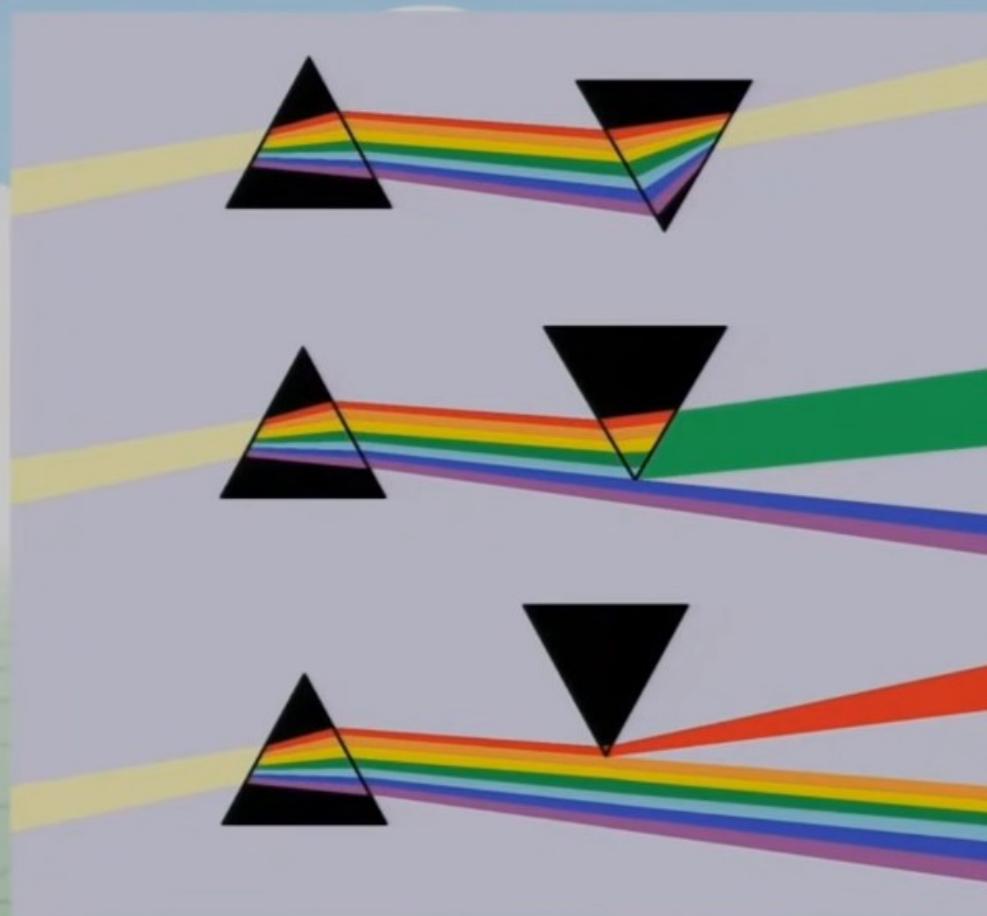
- Like a chord in music
- Can be described as a sum of its parts
- Relative strength of different frequencies



Sources of light are diverse!

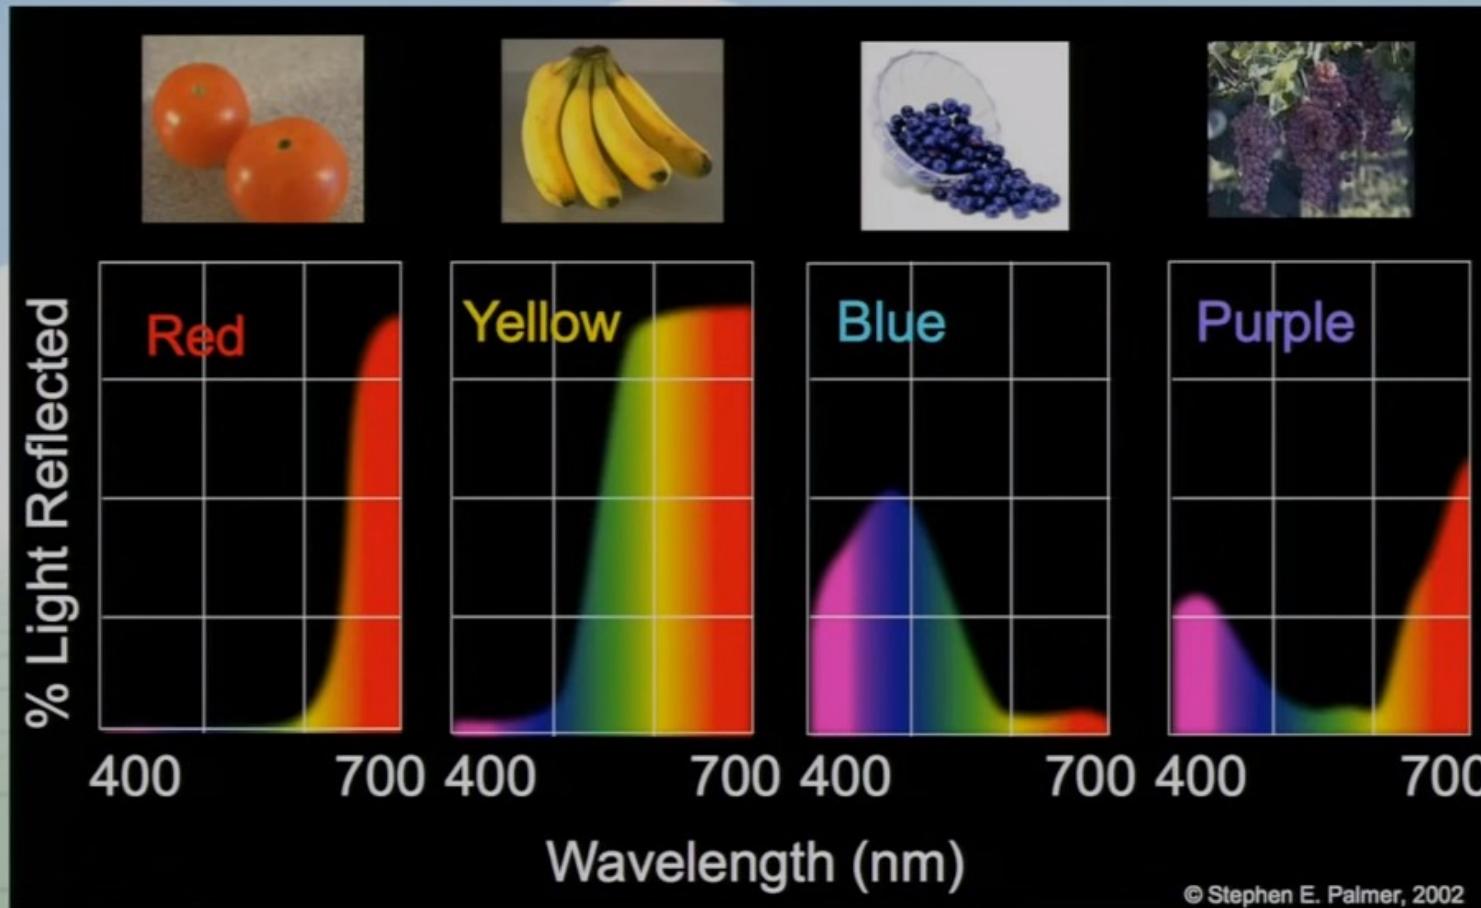


“White” light - all wavelengths



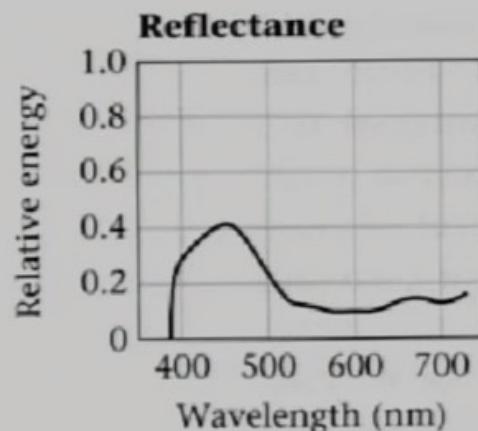
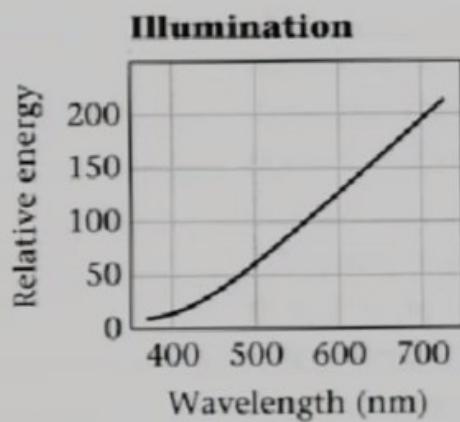
We know this thanks to Newton!

Objects reflect only some light

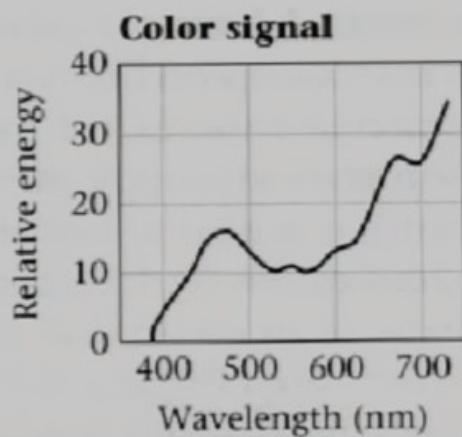


What color is the object?

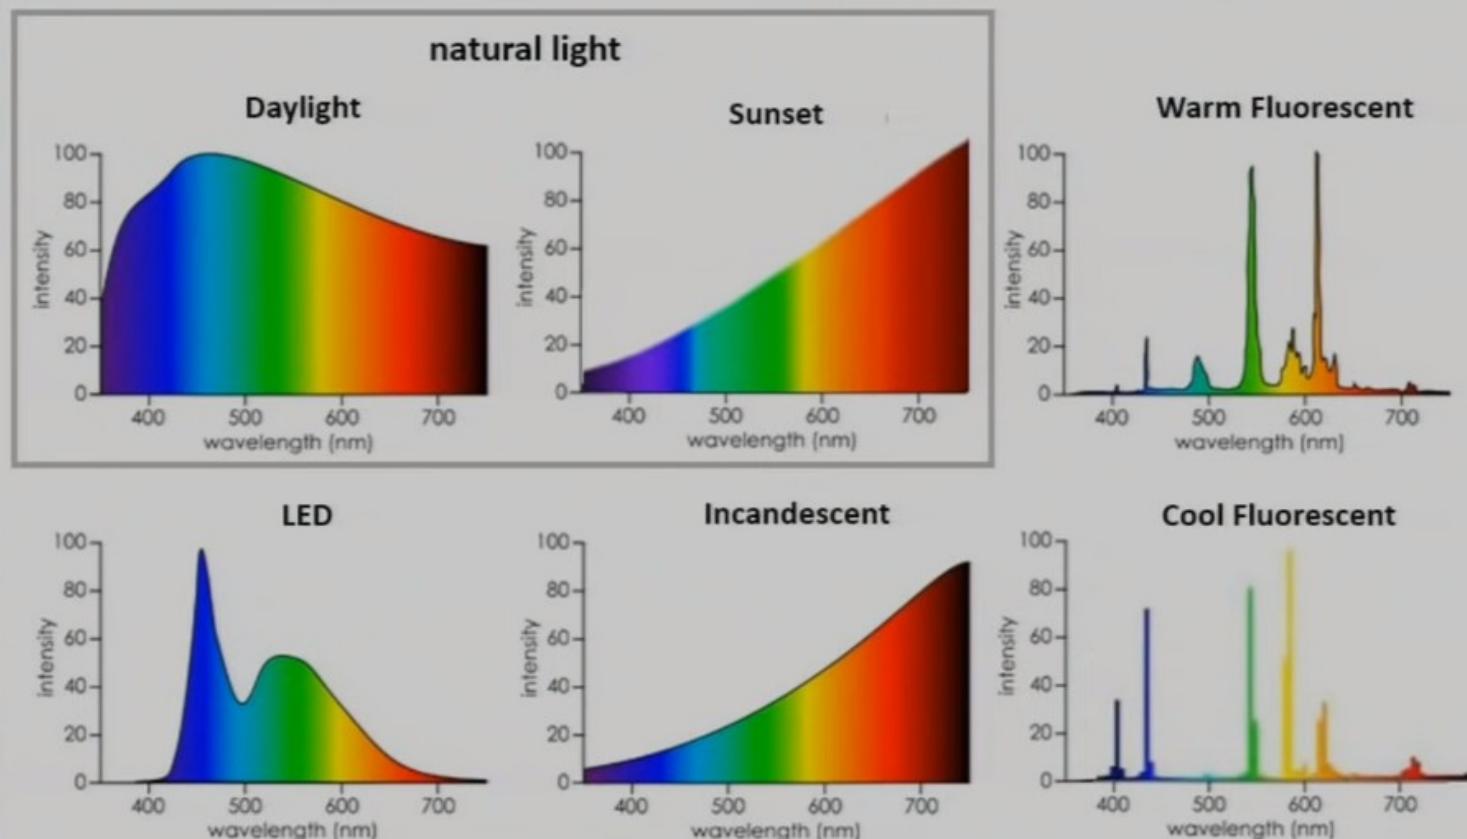
- The “color” of an object depends on both the incident light and the objects reflectance:



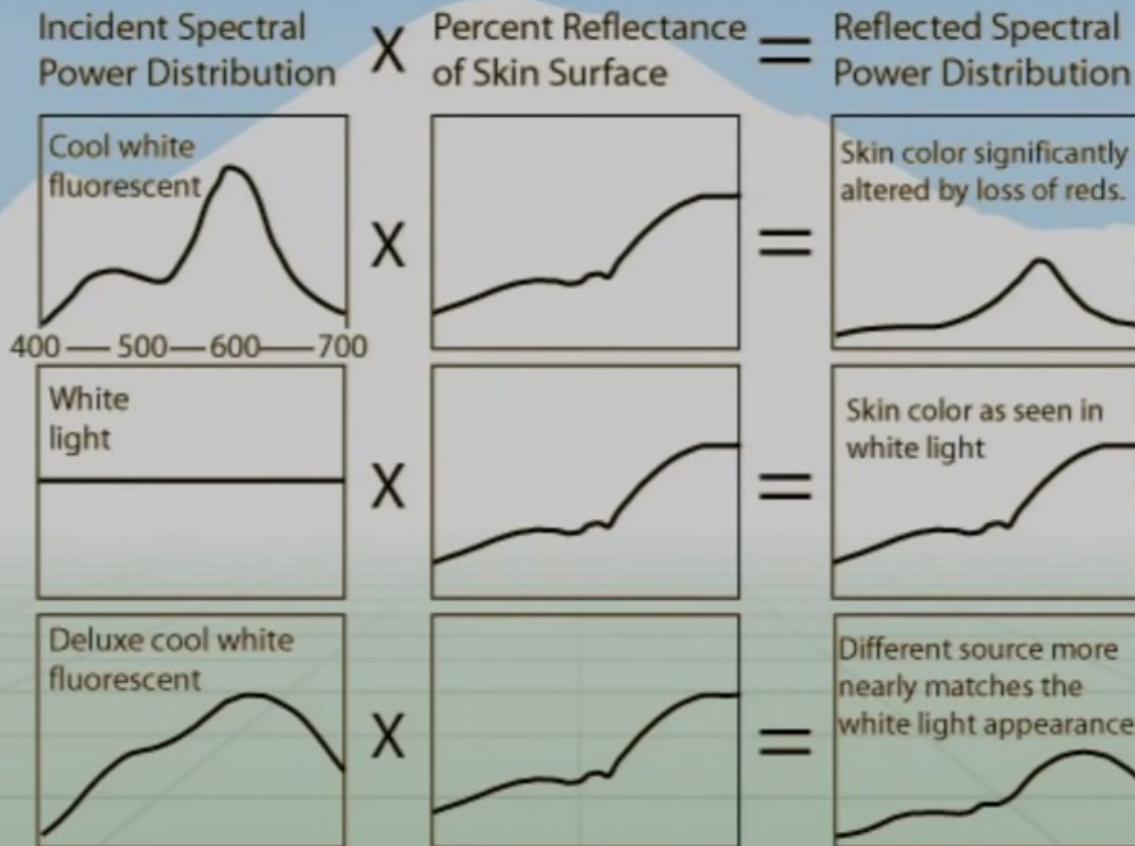
=



Different illumination matters!



Case study: makeup application

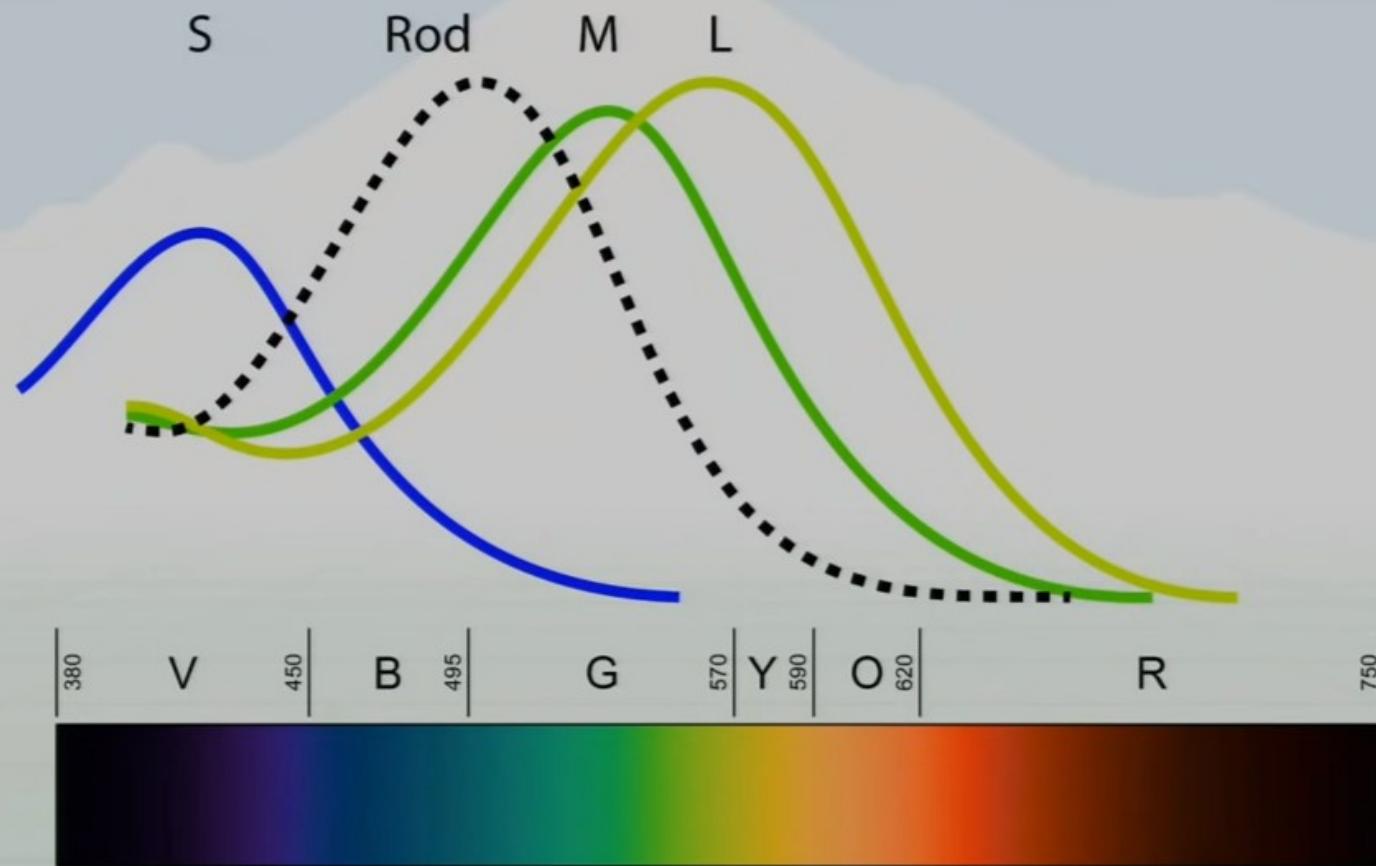


After Williamson and Cummins

Photoreceptors and light

- Each receptor has a responsiveness curve
- Receptors more responsive to some wavelengths, less responsive to others
- Rods: peak around 498 nm
- Cones: 3 kinds
 - Short: peak around 420 nm
 - Medium: peak around 530 nm
 - Long: peak around 560 nm

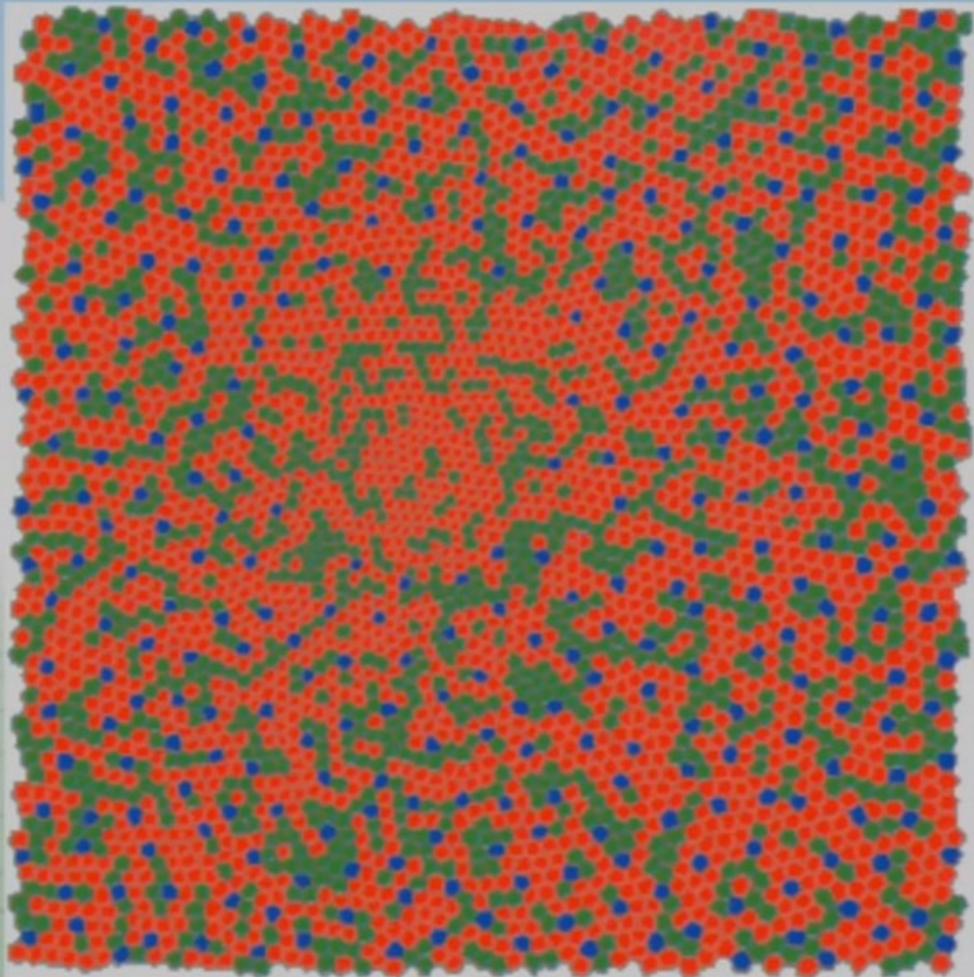
Photoreceptors and light



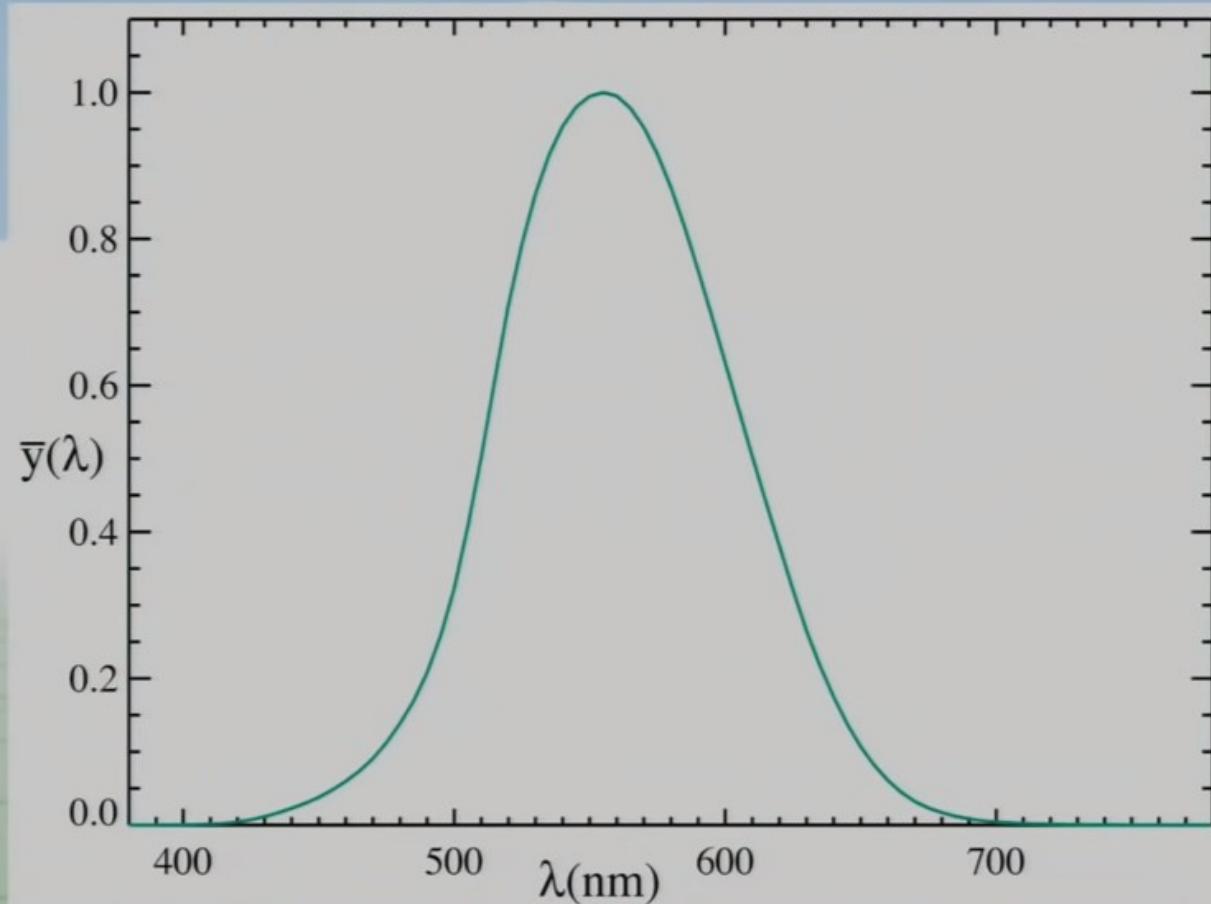
Cones and color

- Our perception of color comes from cones
- Different waveforms provoke different responses
- Each cone has essentially one “output”
- To calculate:
 - Multiply input waveform by response curve
 - Integrate area under the curve
- The “color” we see is the relative activation of the 3 kinds of cones

All cones are not equal



Different wavelengths are brighter

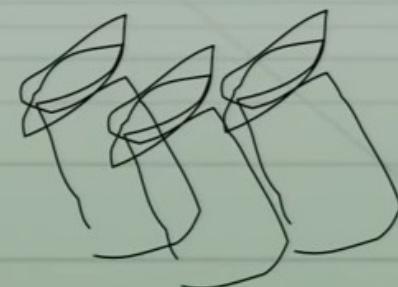
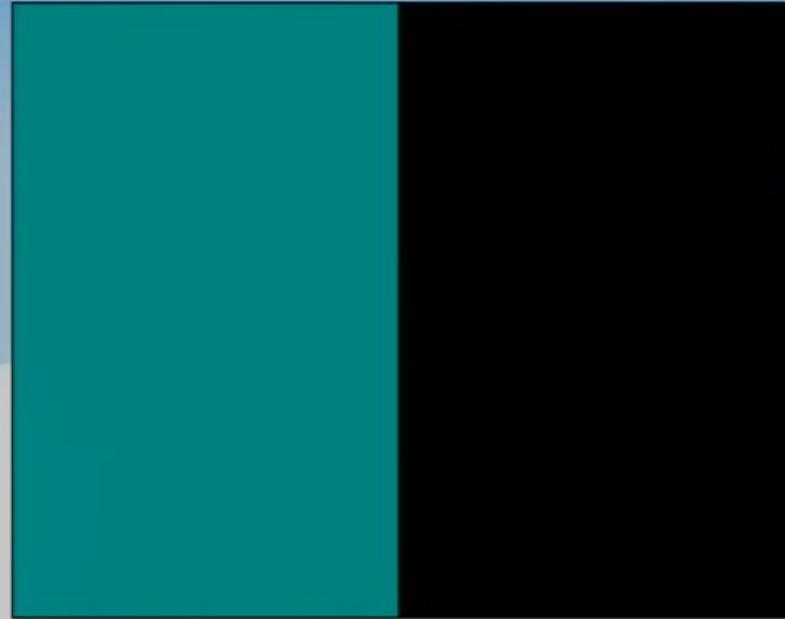


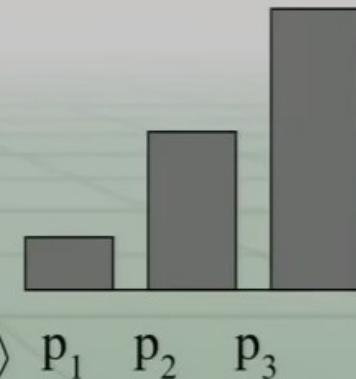
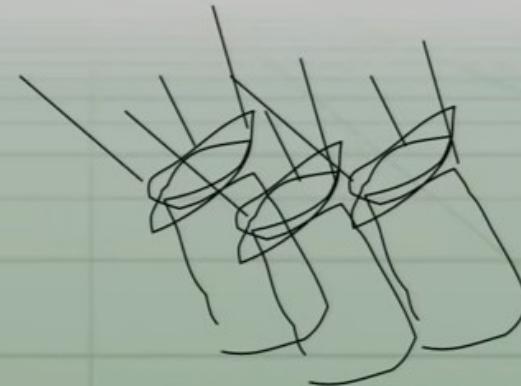
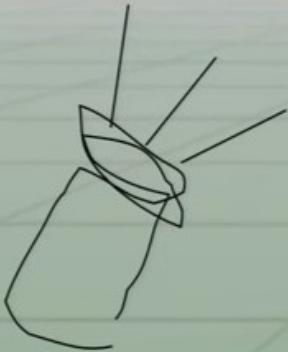
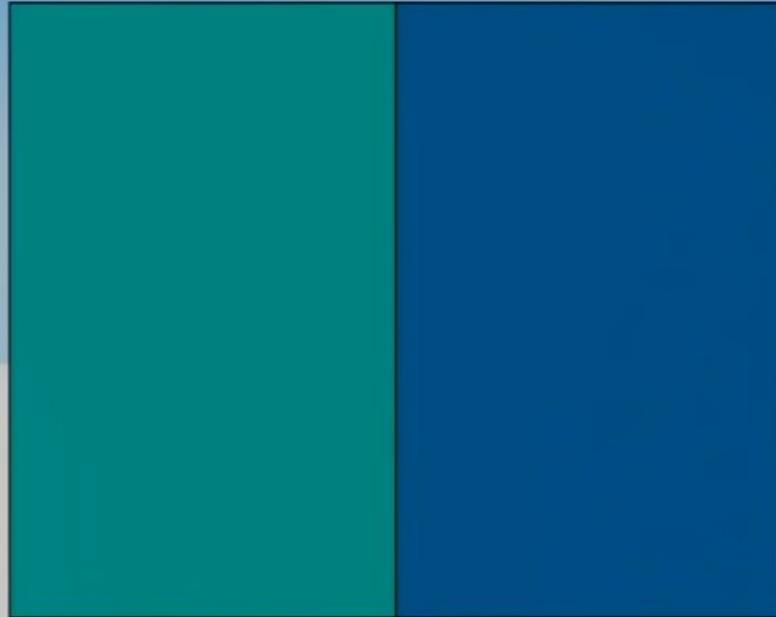
This is hard to read

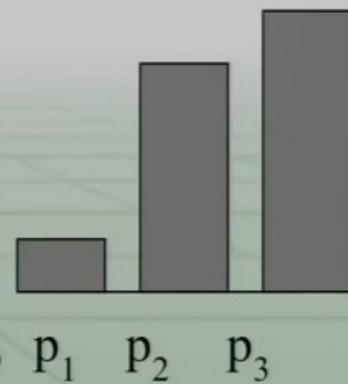
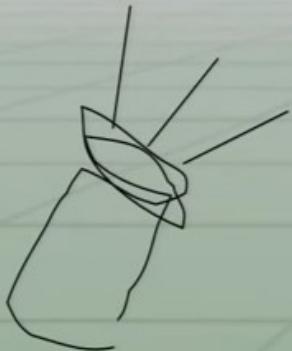
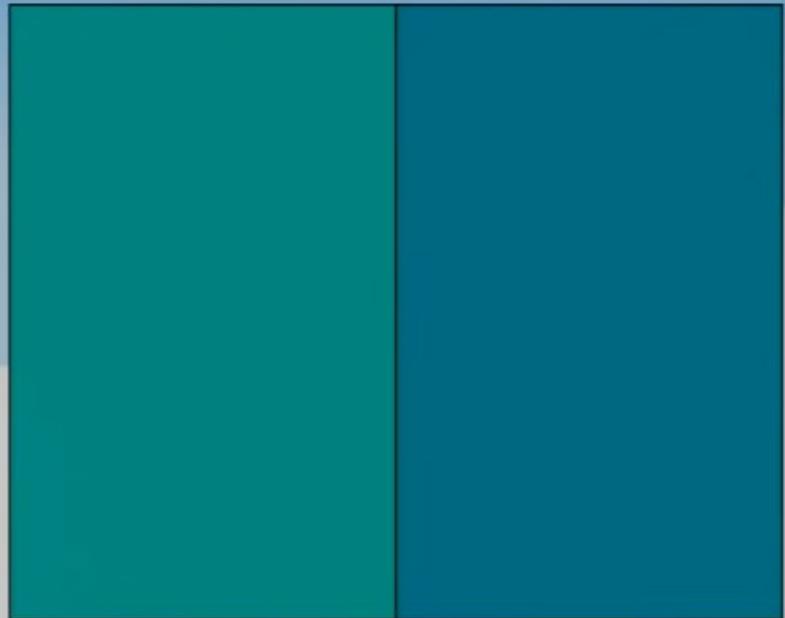
This is easier to read

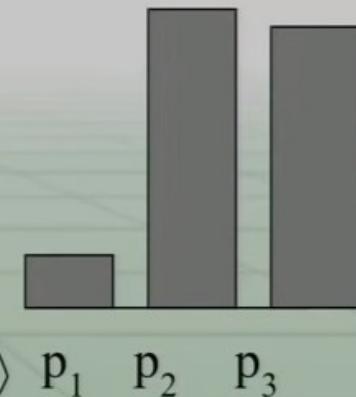
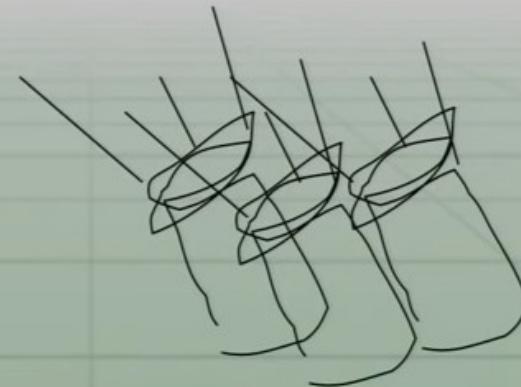
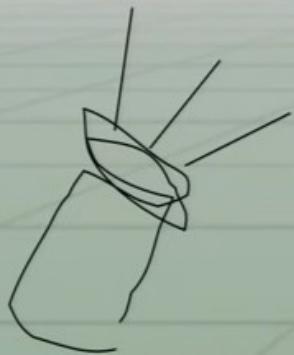
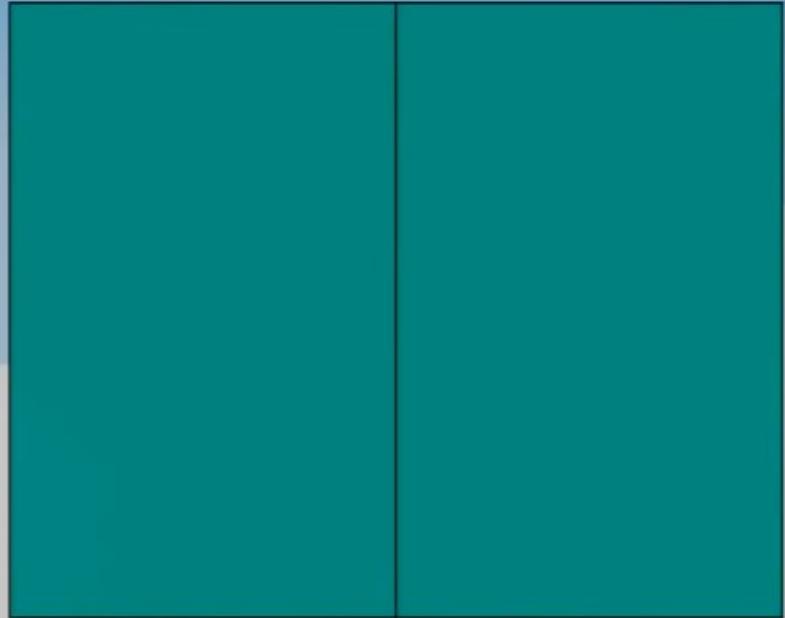
Many variations, what do they see?

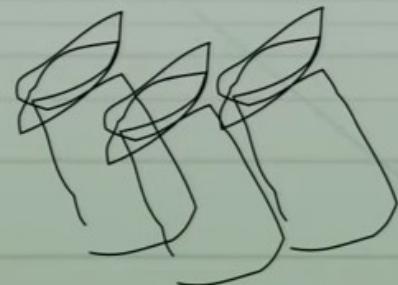
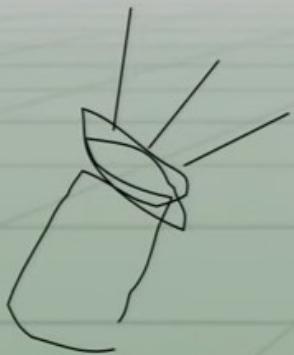
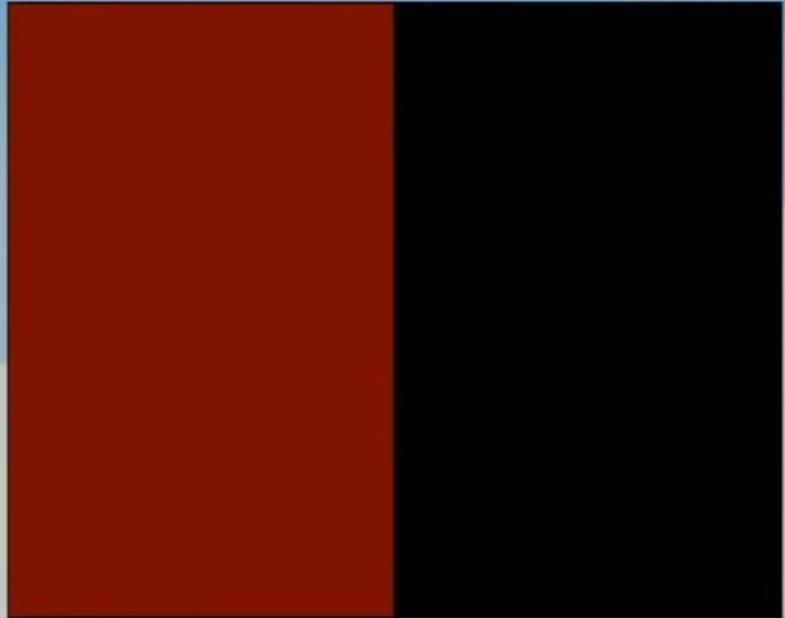
| State | Types of cone cells | Approx. number of colors perceived | Carriers |
|---------------|---------------------|------------------------------------|---|
| Monochromacy | 1 | 100 | marine mammals, owl monkey, Australian sea lion, achromat primates |
| Dichromacy | 2 | 10,000 | most terrestrial non-primate mammals, color blind primates |
| Trichromacy | 3 | 1 million | most primates, especially great apes (such as humans), marsupials, some insects (such as honeybees) |
| Tetrachromacy | 4 | 100 million | most reptiles, amphibians, birds and insects, rarely humans |
| Pentachromacy | 5 | 10 billion | some insects (specific species of butterflies), some birds (pigeons for instance) |

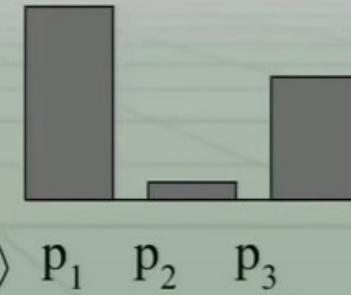
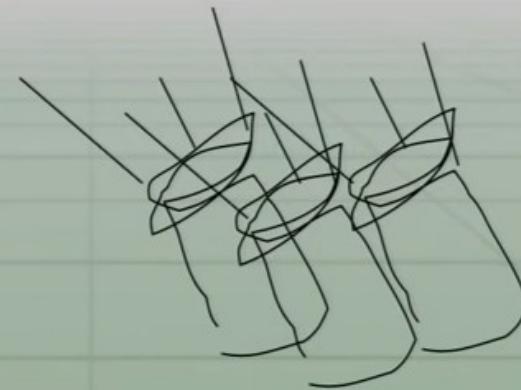
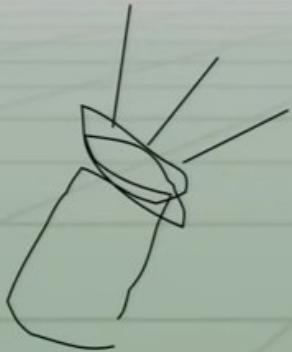
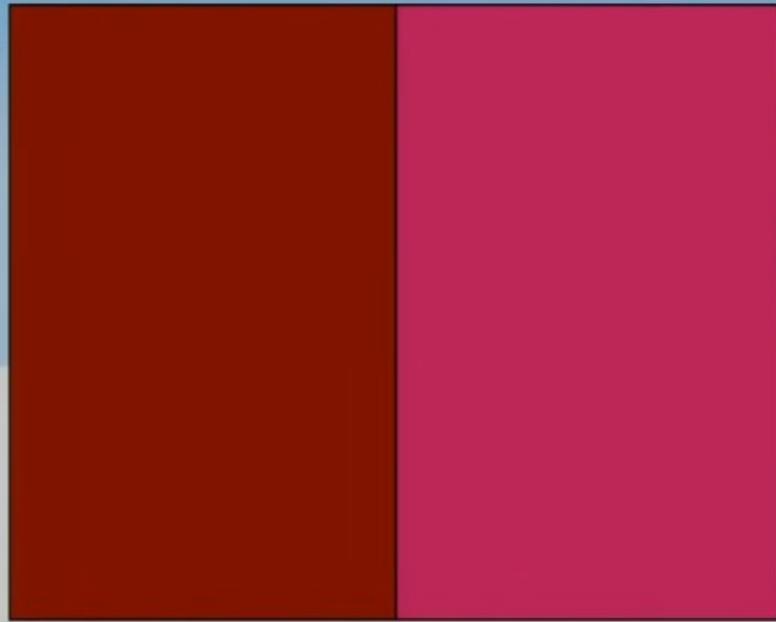


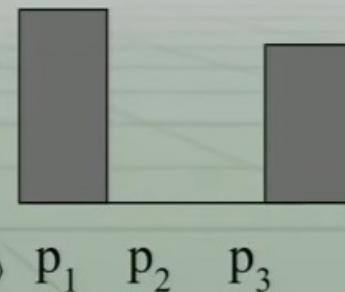
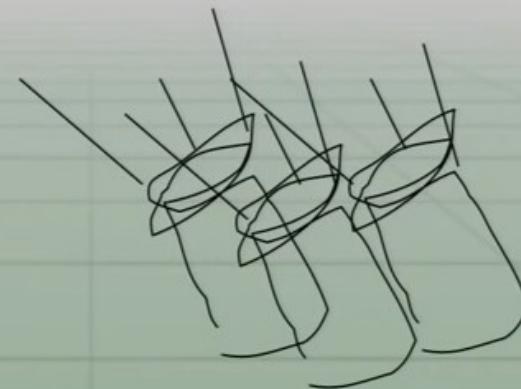
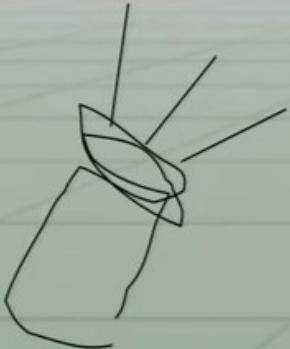
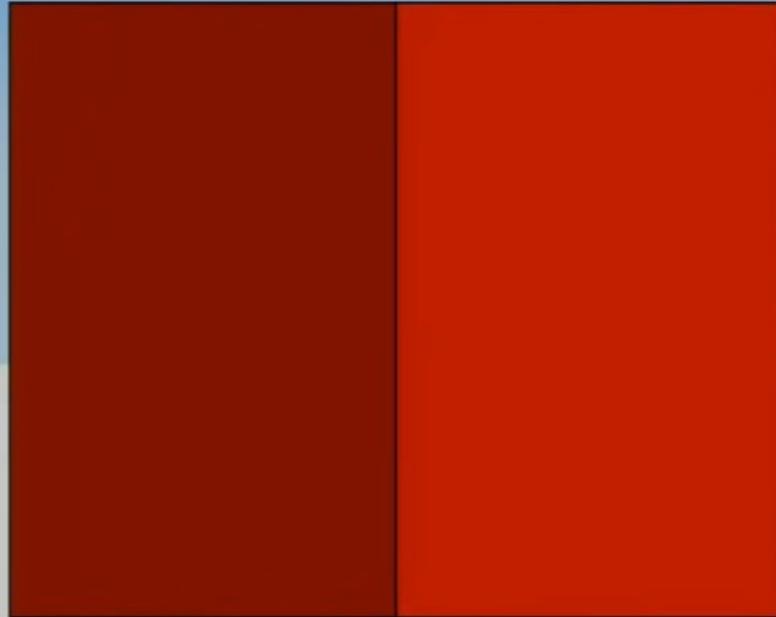


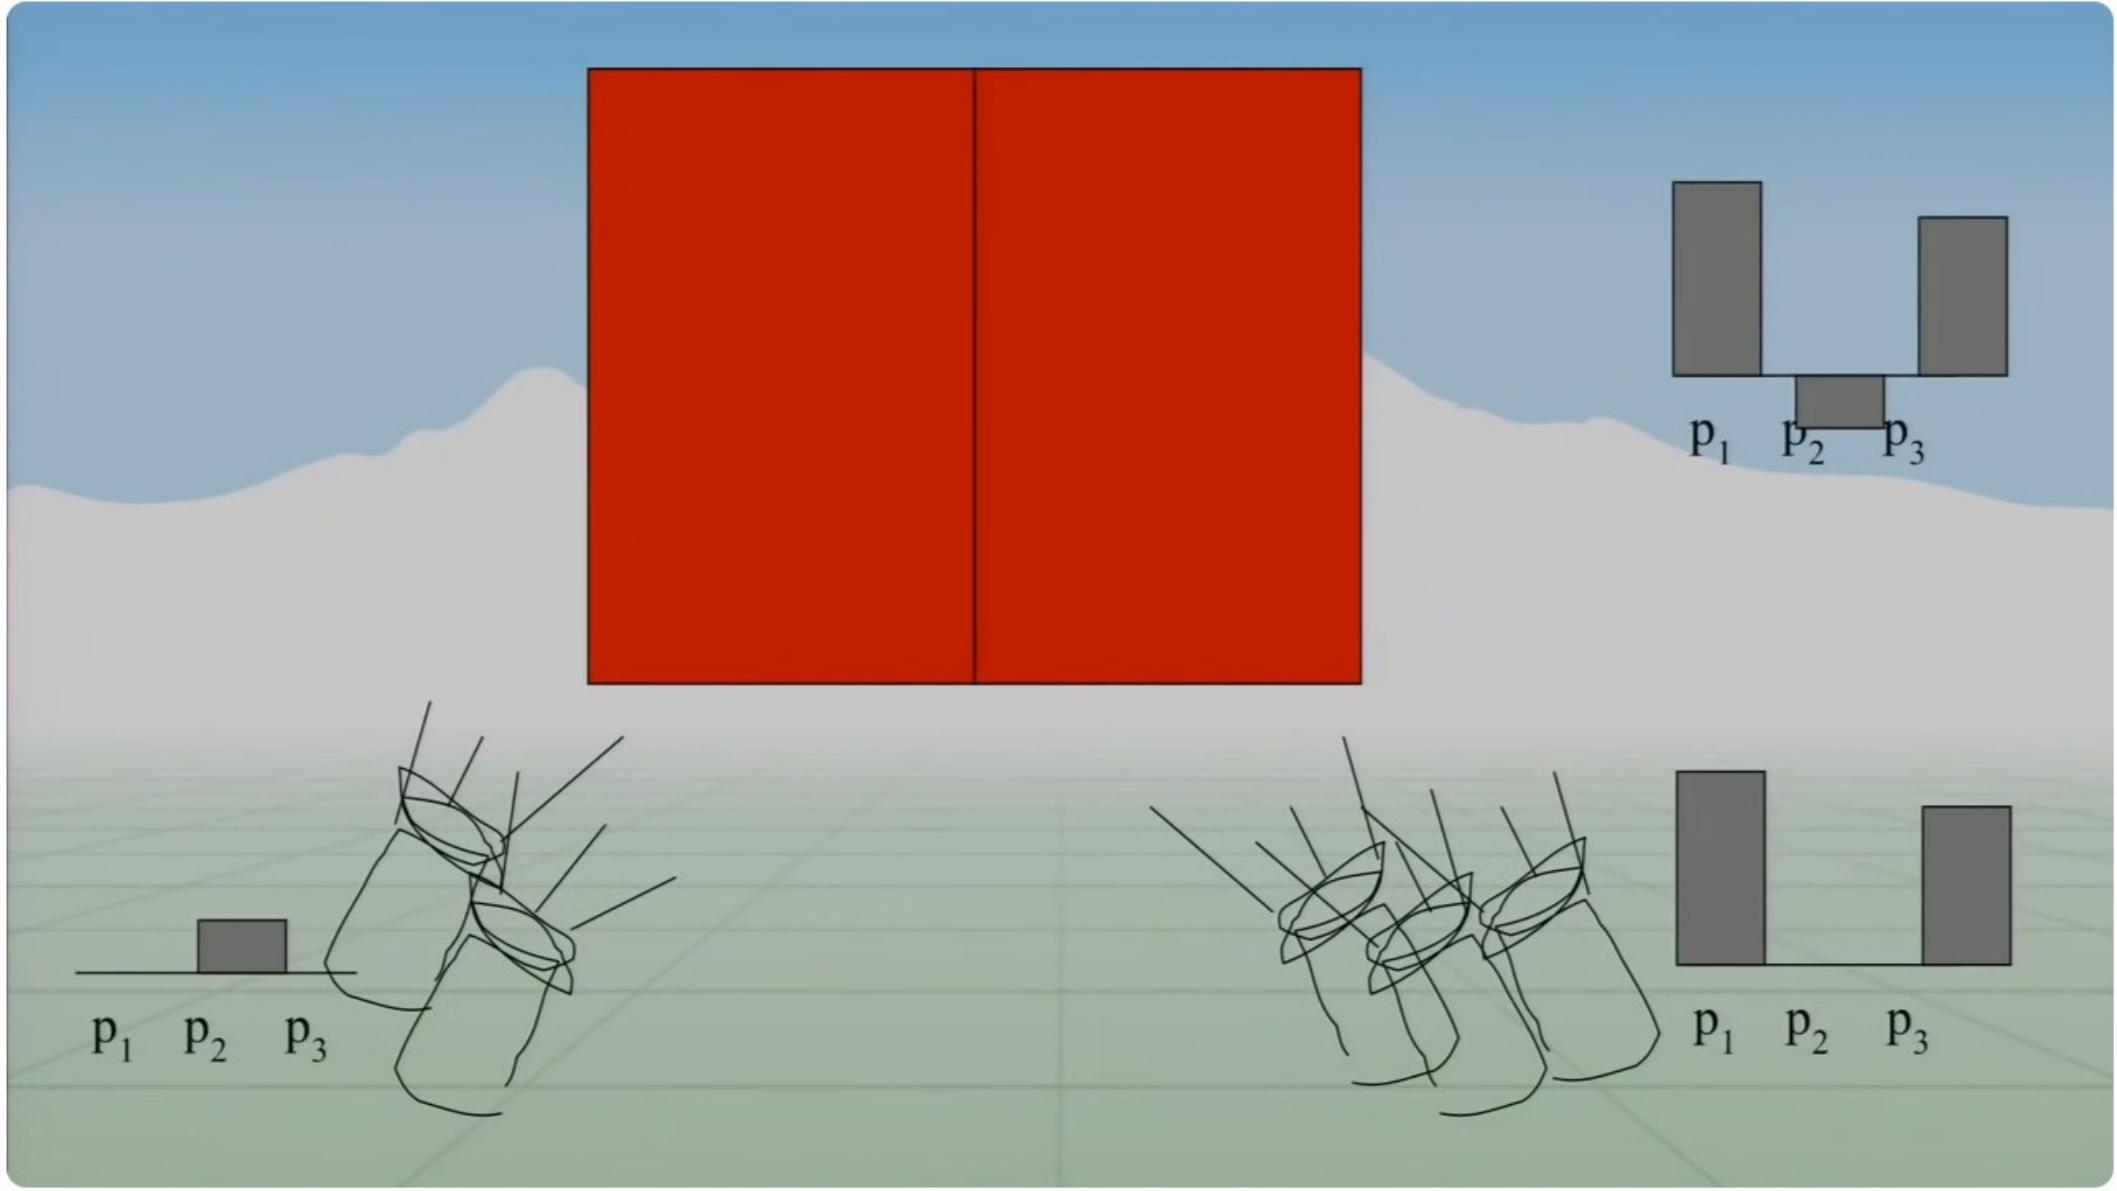


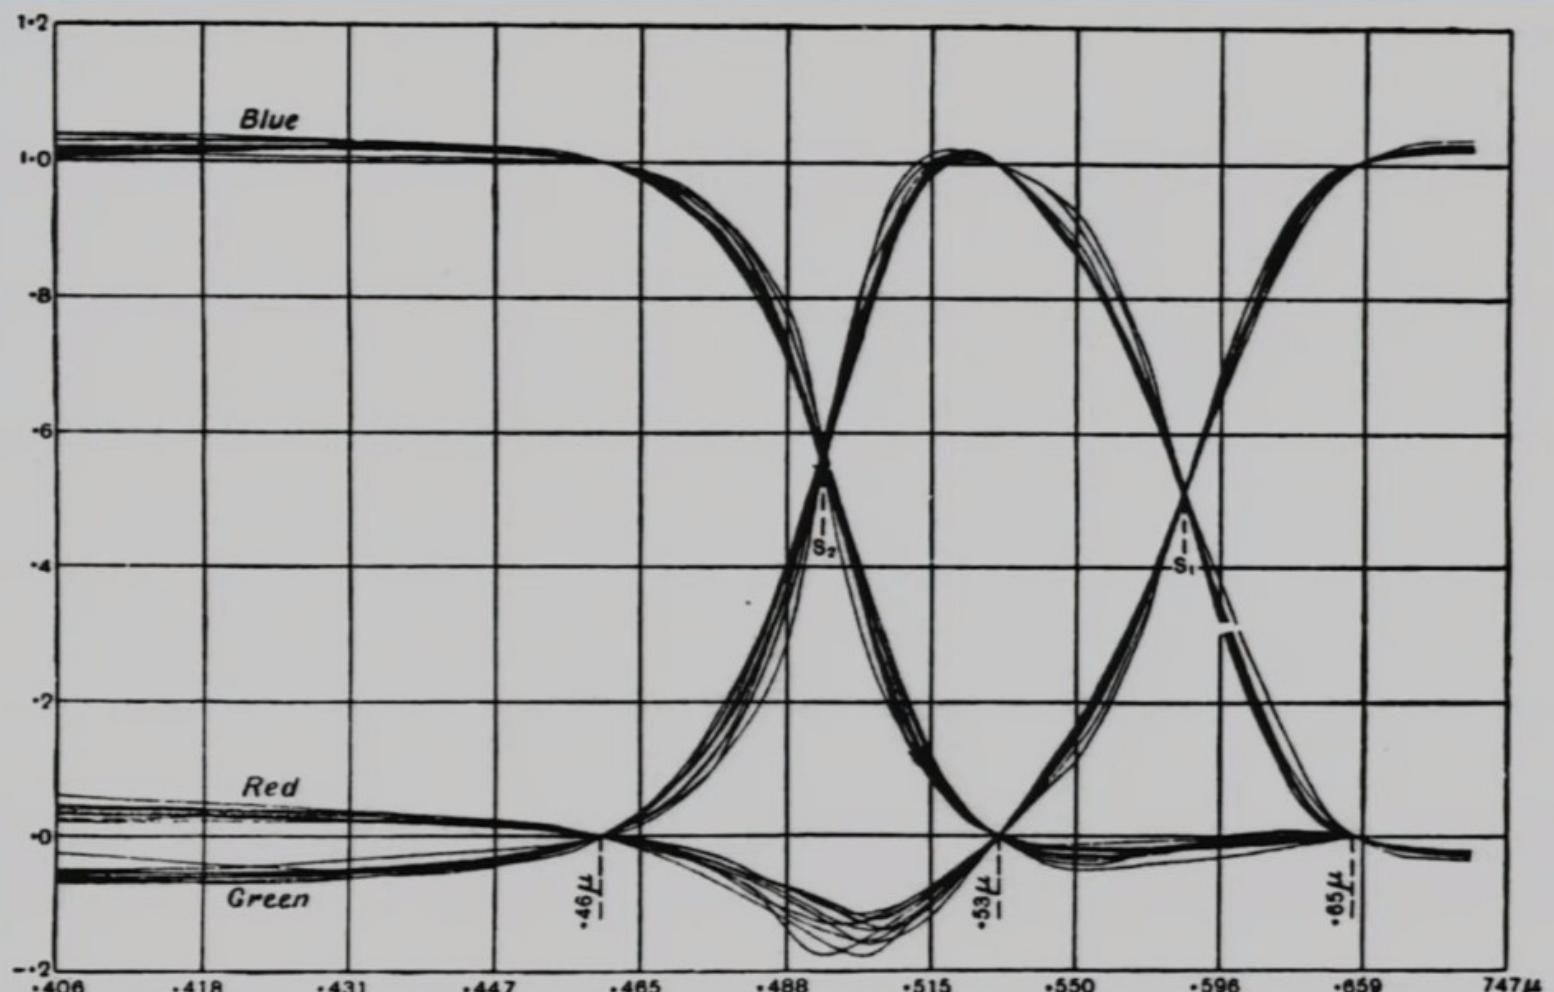


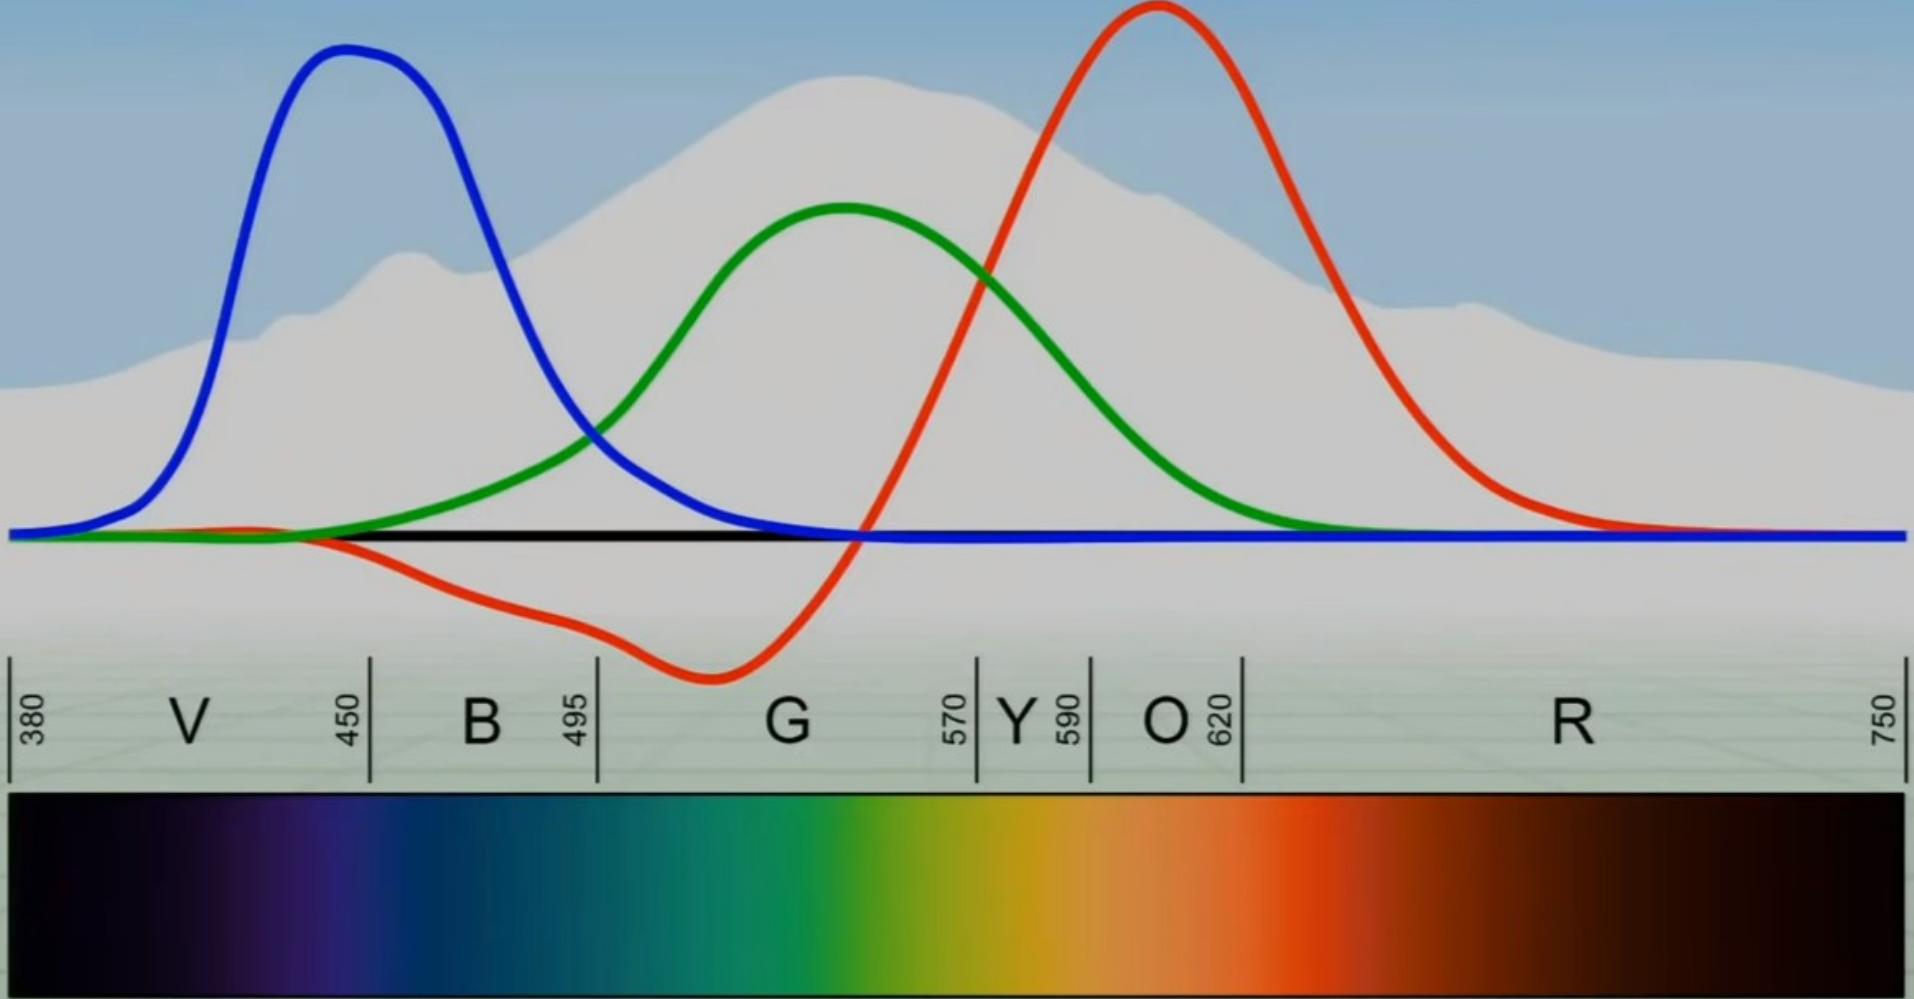


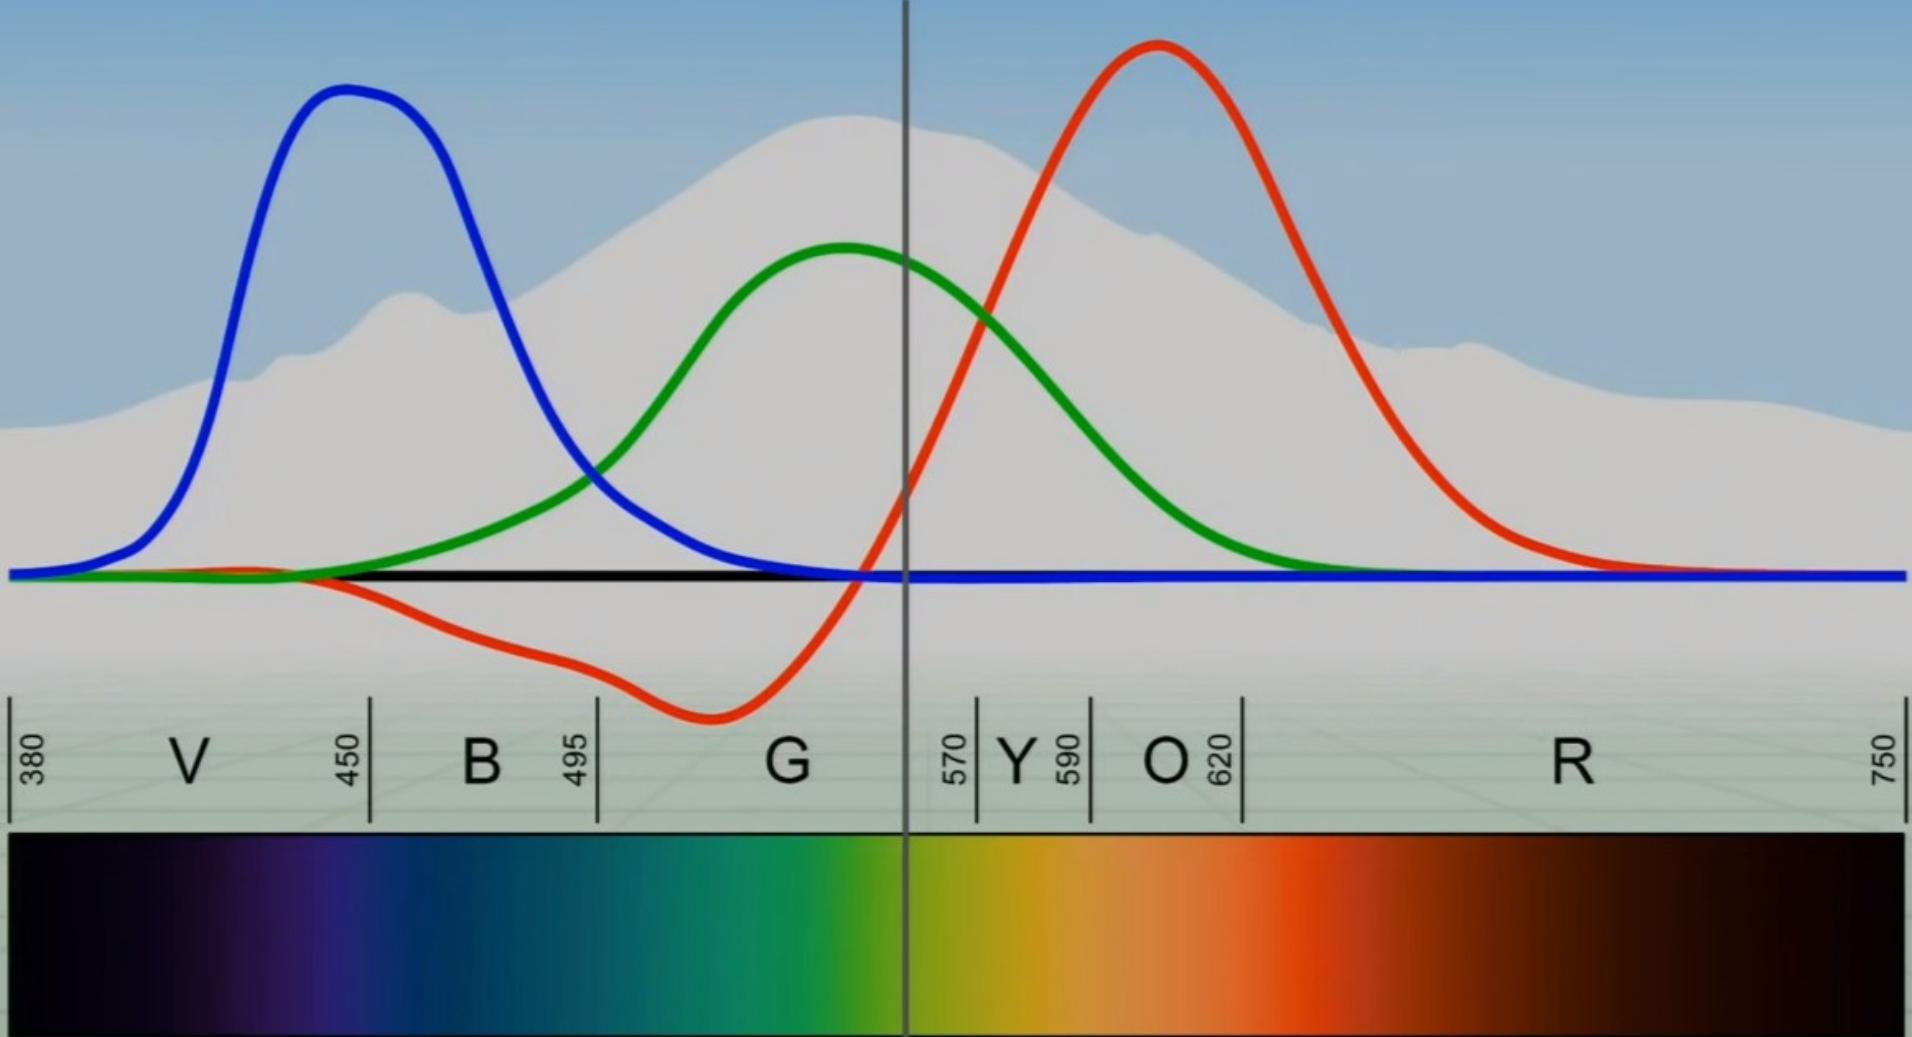


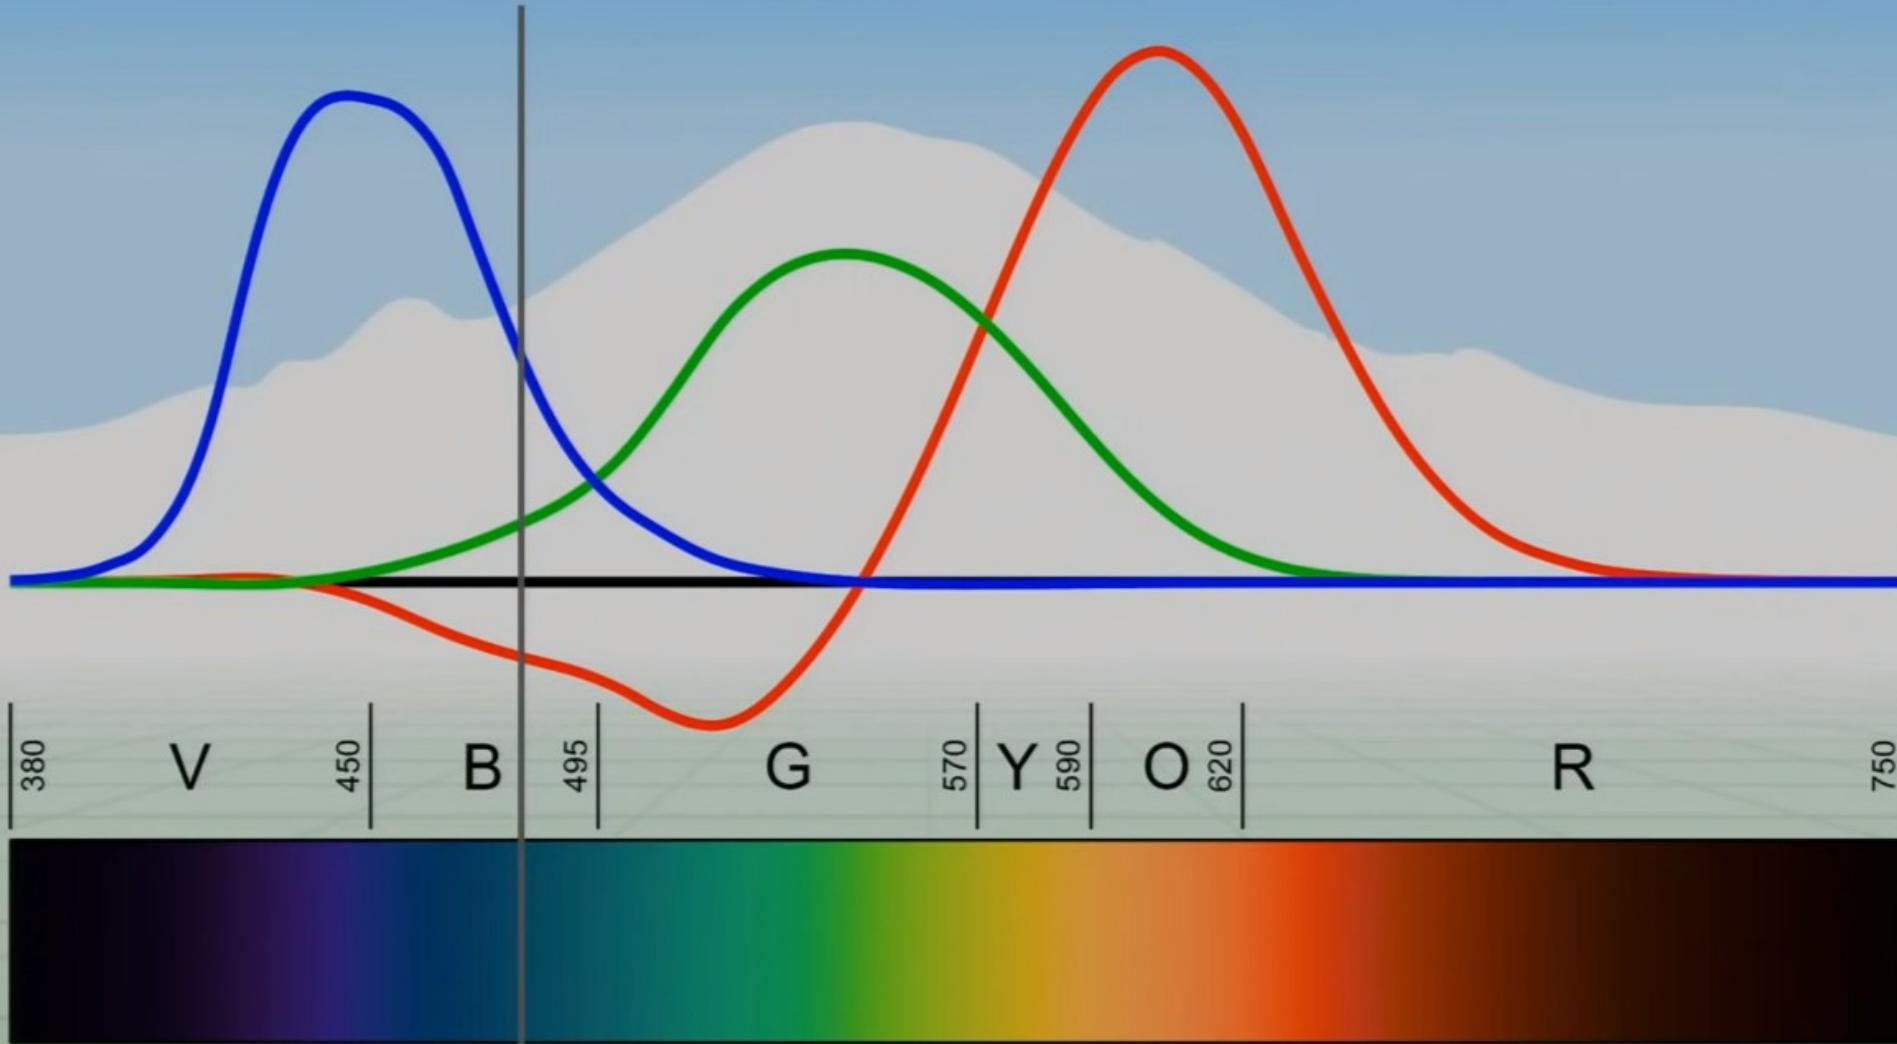








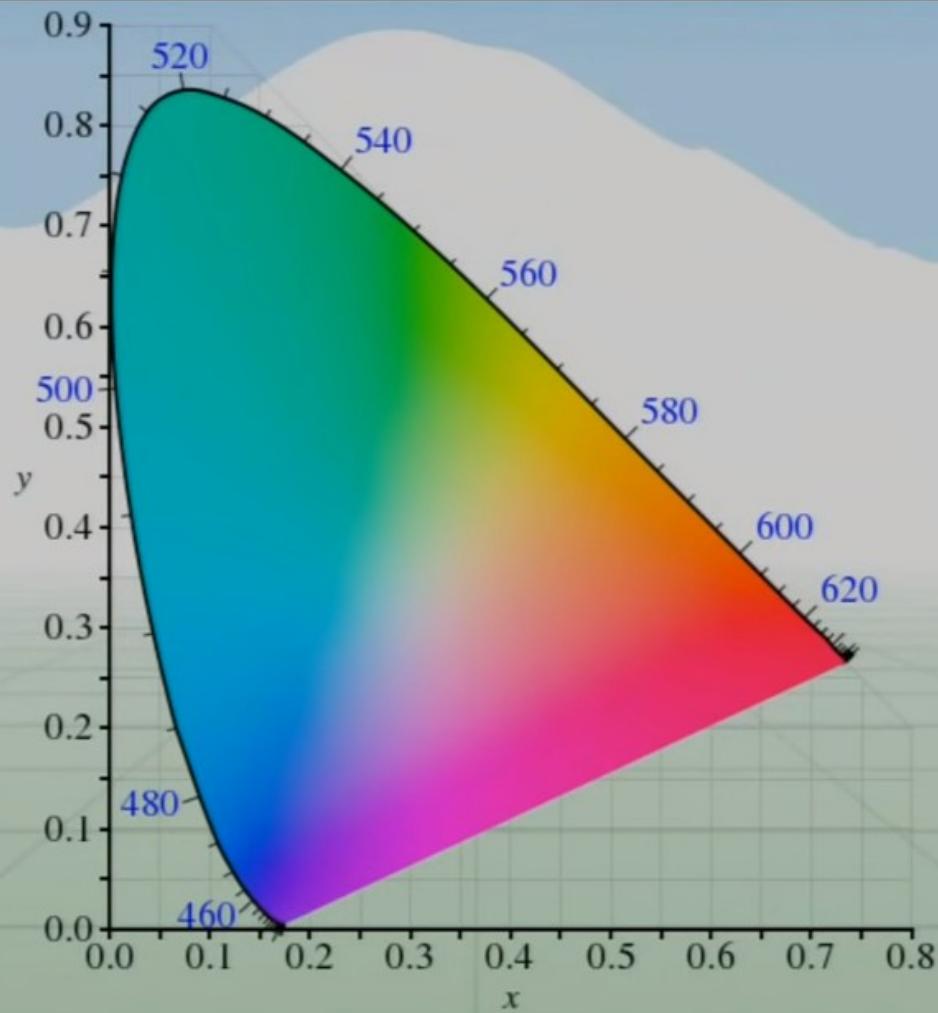




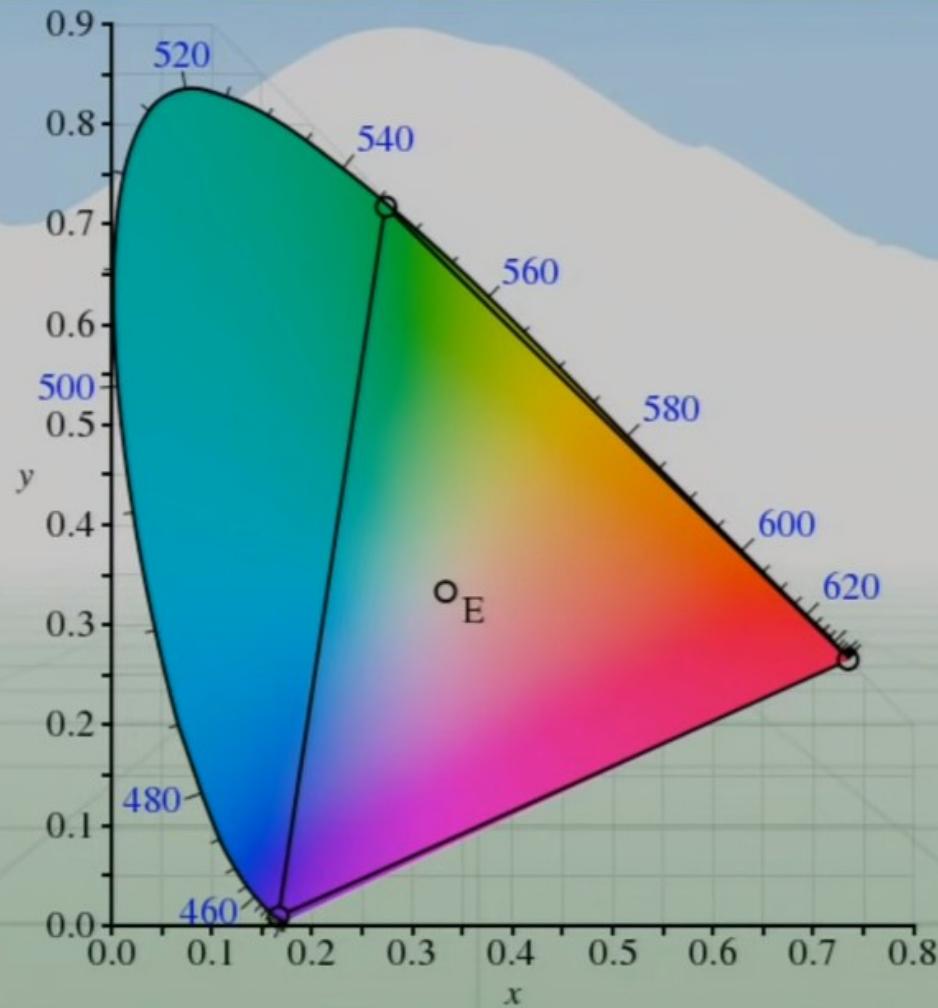
Results:

- Given 3 primaries people can match any color
- People select very similar distributions for a given color
- Colors seem to follow nice, linear rules:
Grassman's laws!
 - $A=B+C \Rightarrow A+D=B+C+D$
 - $A=B+C \Rightarrow nA=nB+nC$
 - $A=B+C$ and $D=B+C \Rightarrow A=D$
- Light is combinations of individual wavelengths
 - If we can match any wavelength we can match any light

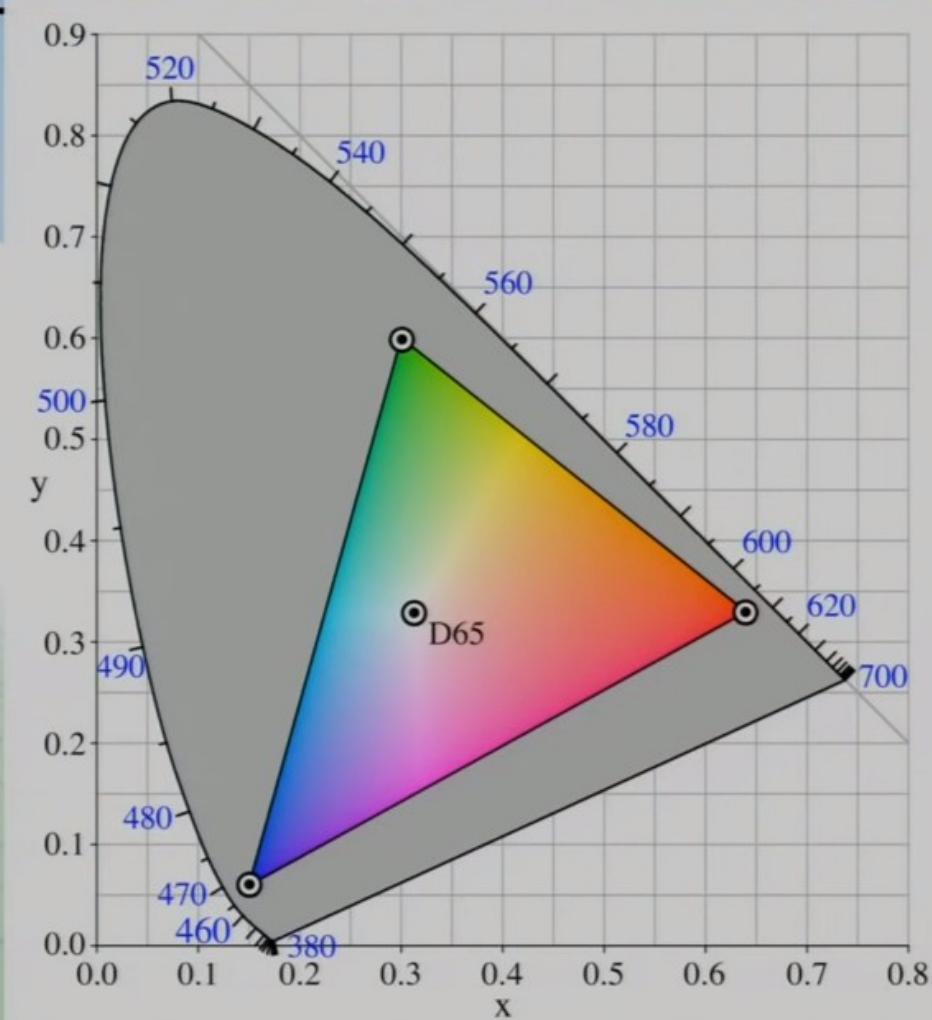
Now we can make a map of color



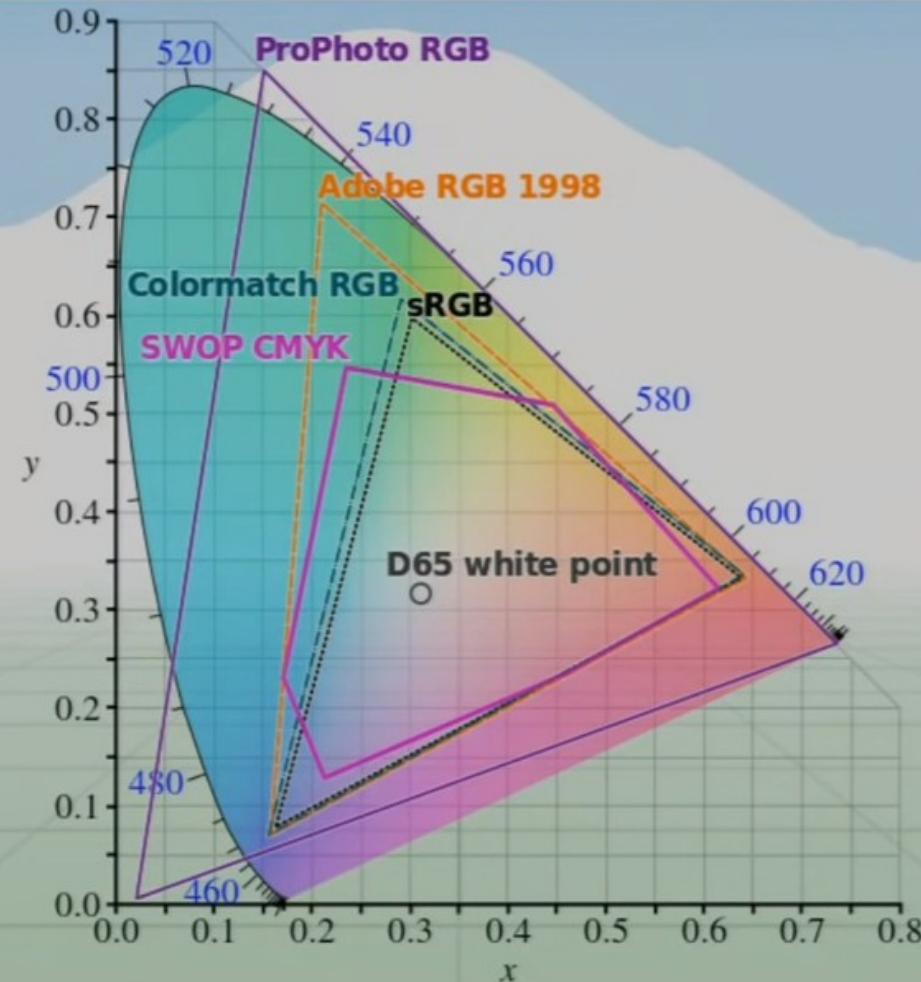
“Theoretical” CIE RGB primaries



Practical sRGB primaries, MSFT 1996



MANY different colorspaces

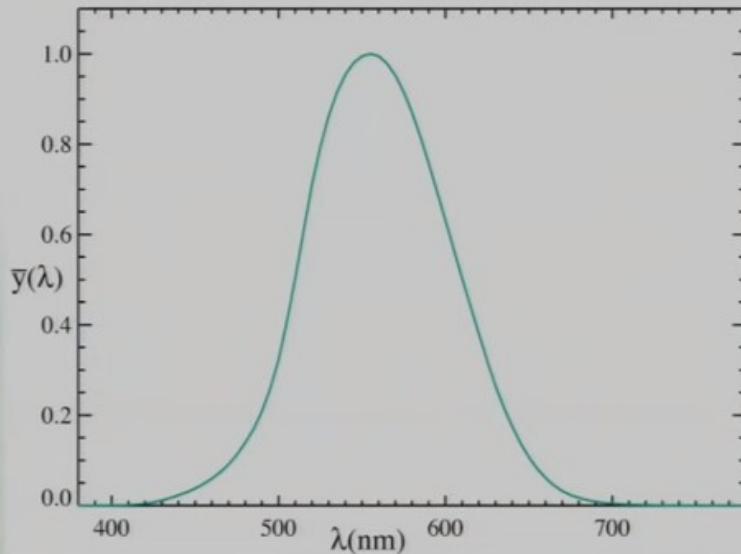


What does this mean for computers?

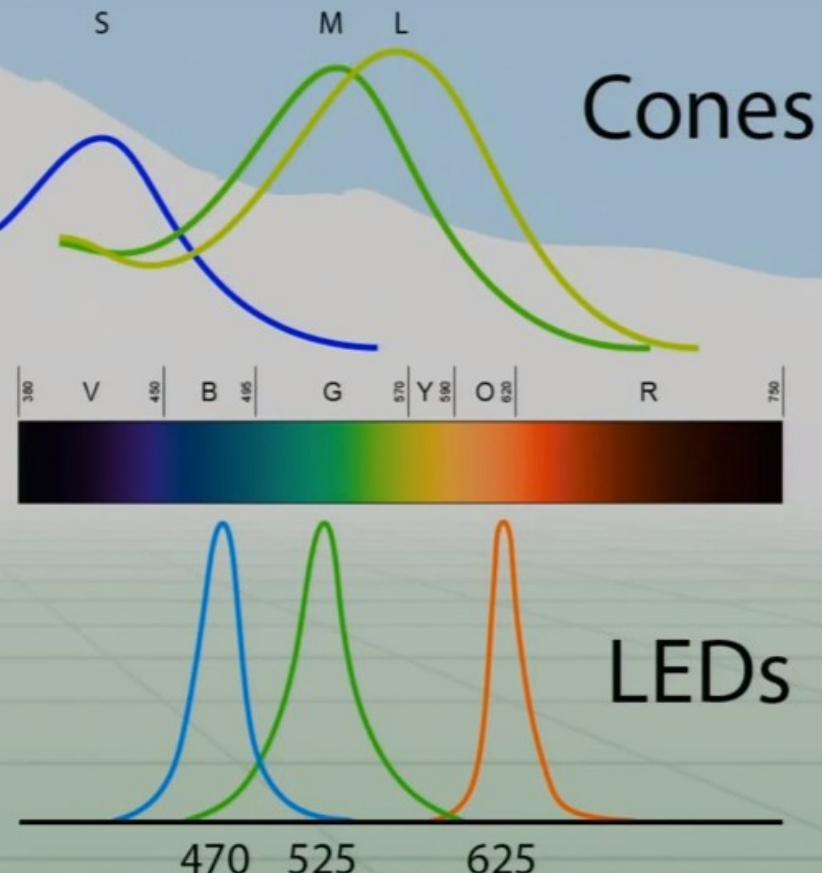
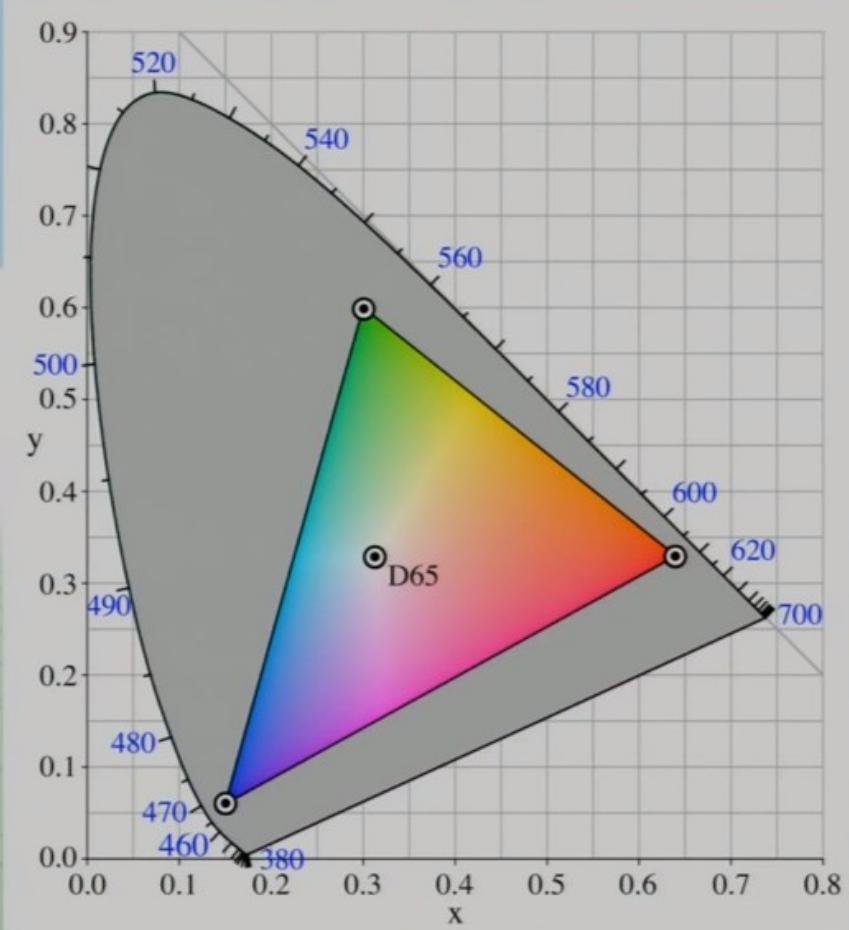
- We represent images as grids of pixels
- Each pixel has a color, 3 components: RGB
- Not every color can be represented in RGB!
 - Have to go out in the real world sometimes
- RGB is made to trick humans, not be accurate
- sRGB is not actually linear, gamma compressed
 - Humans see differences between dark tones more than bright
 - Compress light tones, expand dark tones, more efficient
- Can represent color with 3 numbers
 - #ff00ff; (1.0, 0.0, 1.0); 255,0,255; etc....

Grayscale - making color images not

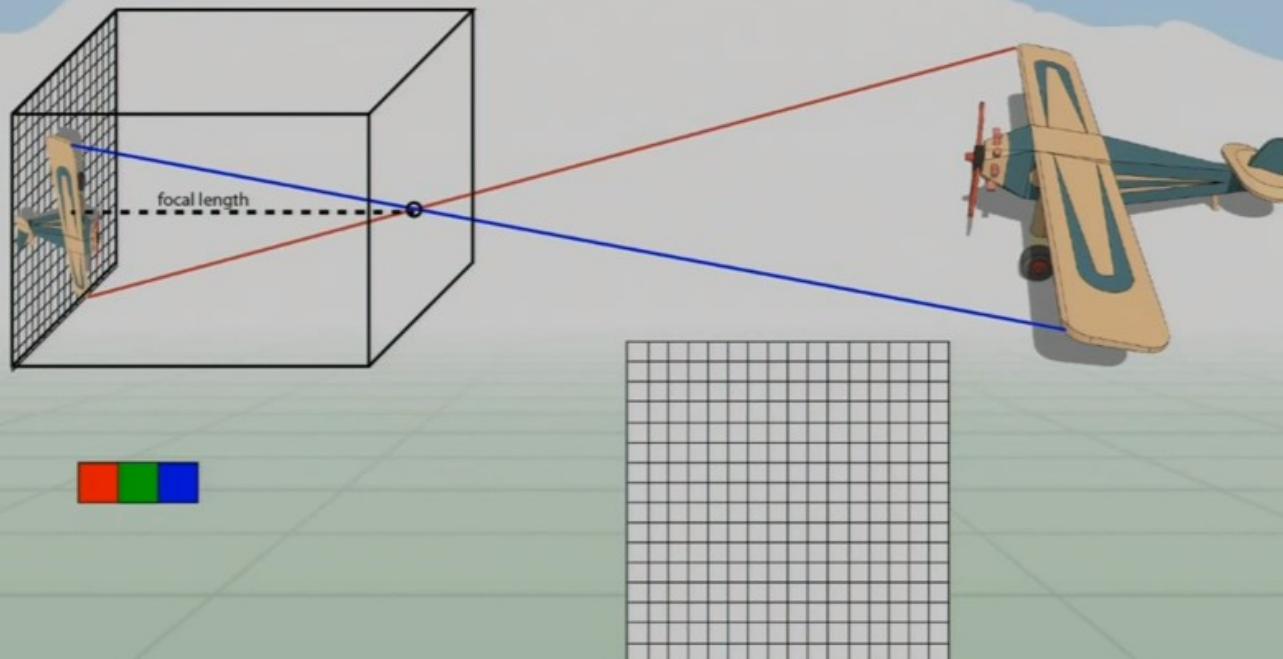
- We can simulate monochromatic images from RGB
- Want a good approximation of how “bright” the image is without color information
- $(R+B+G/3)$ - looks weird
- We *should*
 - Gamma decompress
 - Calculate lightness
 - Gamma compress
- We can just operate on sRGB
 - Typically $\sim .30R + .59G + .11B$



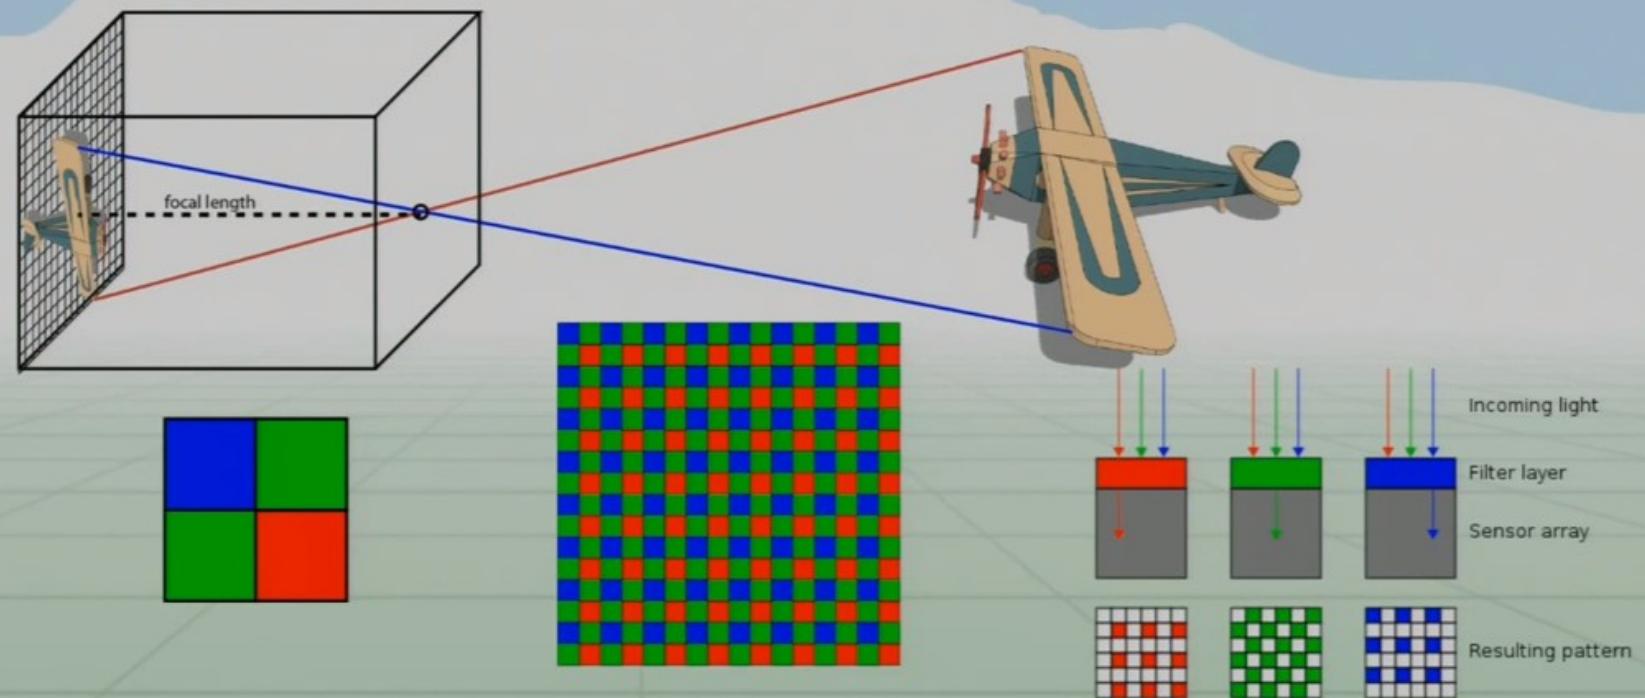
sRGB: most widely used colorspace



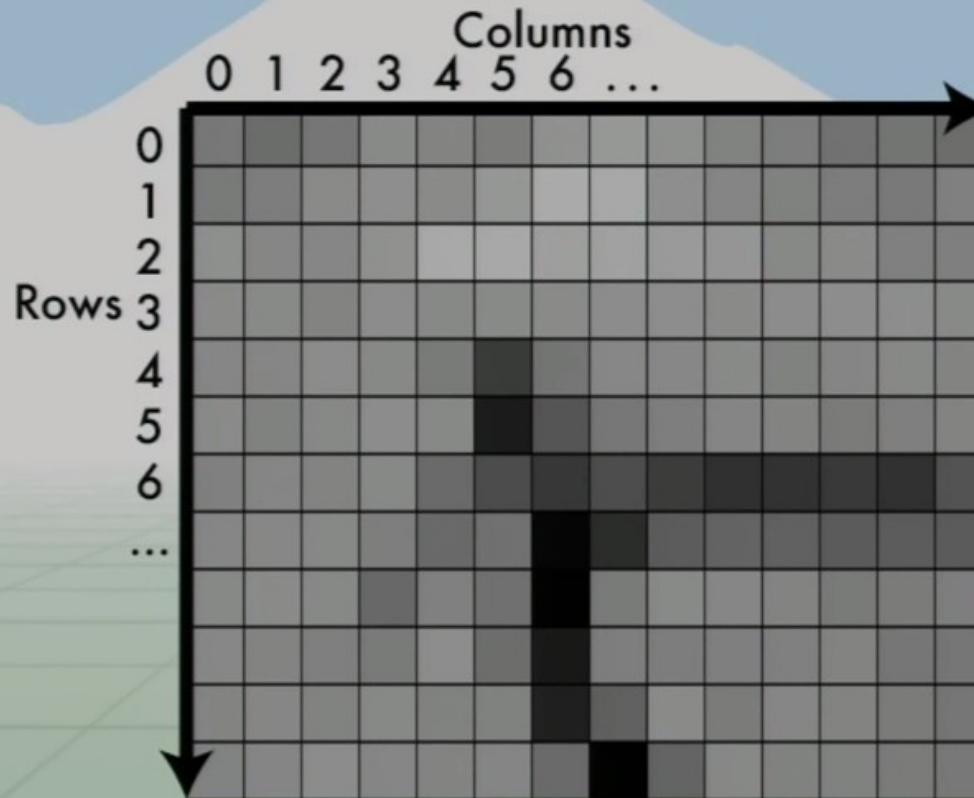
How do we record color?



Bayer pattern for CMOS sensors



An image is a matrix of light



Values in matrix = how much light

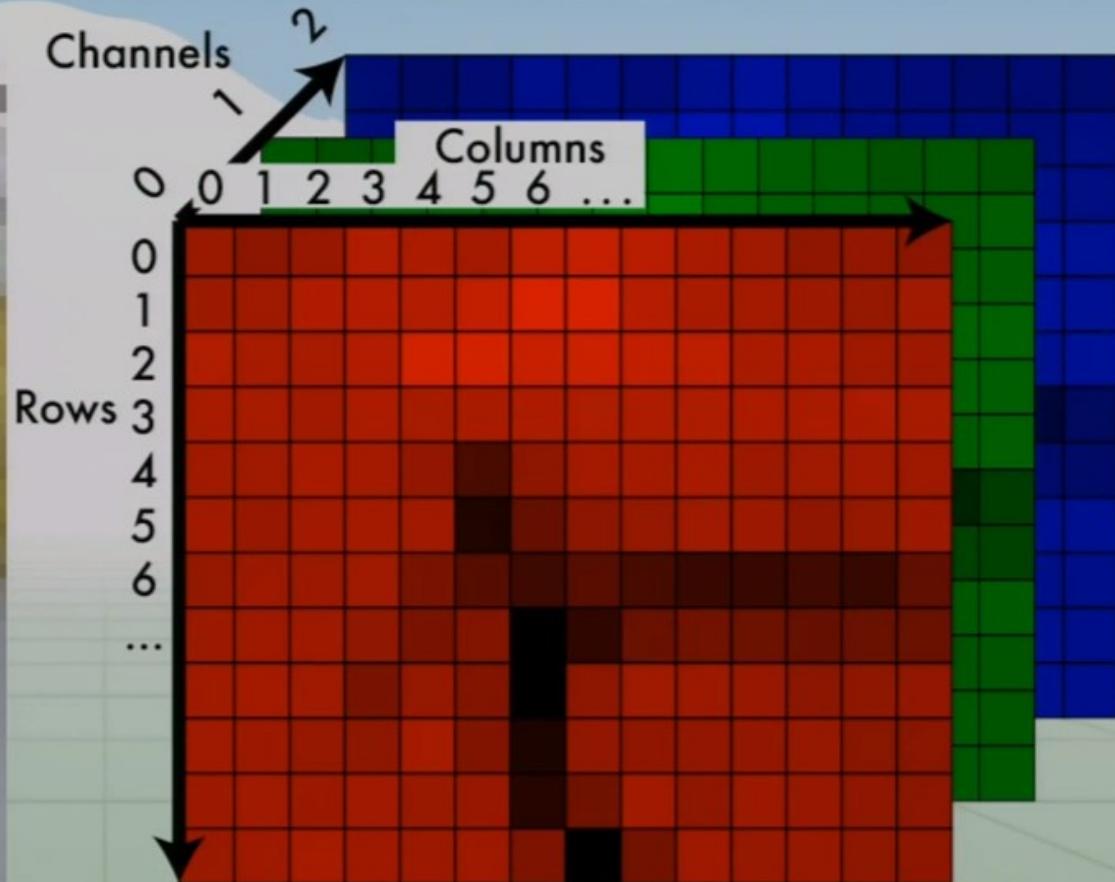
| | | Columns | | | | | | | | | | | | | |
|------|-----|---------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| | | 0 | 1 | 2 | 3 | 4 | 5 | 6 | ... | | | | | | |
| Rows | 0 | 100 | 102 | 107 | 102 | 132 | 146 | 136 | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | 1 | 100 | 102 | 107 | 102 | 132 | 146 | 136 | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | 2 | 100 | 102 | 107 | 102 | 132 | 146 | 136 | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | 3 | 100 | 102 | 107 | 102 | 132 | 146 | 136 | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | 4 | 100 | 102 | 107 | 102 | 132 | 146 | 136 | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | 5 | 100 | 102 | 107 | 102 | 132 | 30 | 60 | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | 6 | 100 | 102 | 107 | 102 | 132 | 40 | 20 | 50 | 32 | 20 | 20 | 24 | 30 | 62 |
| | ... | 100 | 102 | 107 | 102 | 132 | 71 | | 156 | 51 | 57 | 57 | 58 | 62 | 58 |
| | | 100 | 102 | 107 | 102 | 132 | 69 | | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | | 100 | 102 | 107 | 102 | 132 | 89 | 12 | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | | 100 | 102 | 107 | 102 | 132 | 146 | 13 | 45 | 148 | 122 | 115 | 104 | 105 | 103 |
| | | 100 | 102 | 107 | 102 | 132 | 146 | 46 | | 42 | 122 | 115 | 104 | 105 | 103 |

Values in matrix = how much light

- Higher = more light
 - Lower = less light
 - Bounded
 - No light = 0
 - Sensor/device limit = max
 - Typical ranges:
 - [0-255], fit into byte
 - [0-1], floating point
 - Called pixels

| | | Columns | | | | | | | | | | | | | |
|------|-----|---------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| | | 0 | 1 | 2 | 3 | 4 | 5 | 6 | ... | | | | | | |
| Rows | 0 | 100 | 102 | 107 | 102 | 132 | 146 | 136 | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | 1 | 100 | 102 | 107 | 102 | 132 | 146 | 136 | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | 2 | 100 | 102 | 107 | 102 | 132 | 146 | 136 | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | 3 | 100 | 102 | 107 | 102 | 132 | 146 | 136 | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | 4 | 100 | 102 | 107 | 102 | 132 | 146 | 136 | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | 5 | 100 | 102 | 107 | 102 | 132 | 30 | 60 | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | 6 | 100 | 102 | 107 | 102 | 132 | 40 | 20 | 50 | 32 | 20 | 20 | 24 | 30 | 62 |
| | ... | 100 | 102 | 107 | 102 | 132 | 71 | | 156 | 51 | 57 | 57 | 58 | 62 | 58 |
| | | 100 | 102 | 107 | 102 | 132 | 69 | | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | | 100 | 102 | 107 | 102 | 132 | 89 | 12 | 156 | 148 | 122 | 115 | 104 | 105 | 103 |
| | | 100 | 102 | 107 | 102 | 132 | 146 | 13 | 45 | 148 | 122 | 115 | 104 | 105 | 103 |
| | | 100 | 102 | 107 | 102 | 132 | 146 | 46 | | 42 | 122 | 115 | 104 | 105 | 103 |

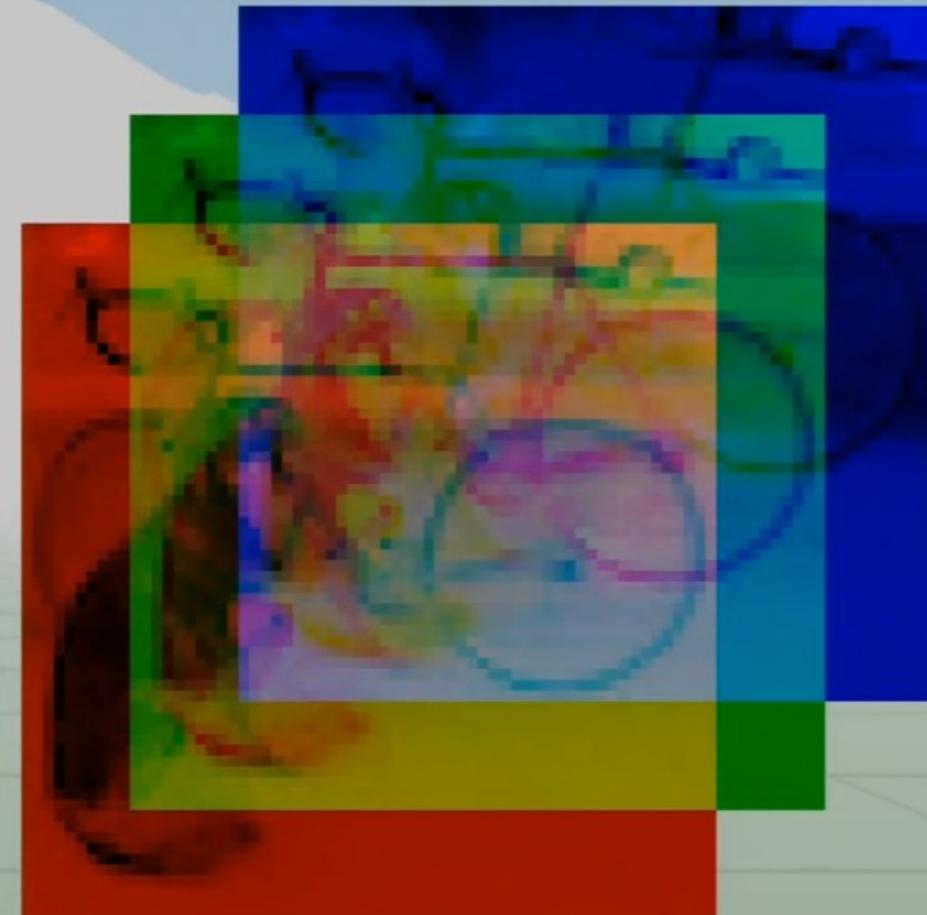
Color image: 3d tensor in colorspace



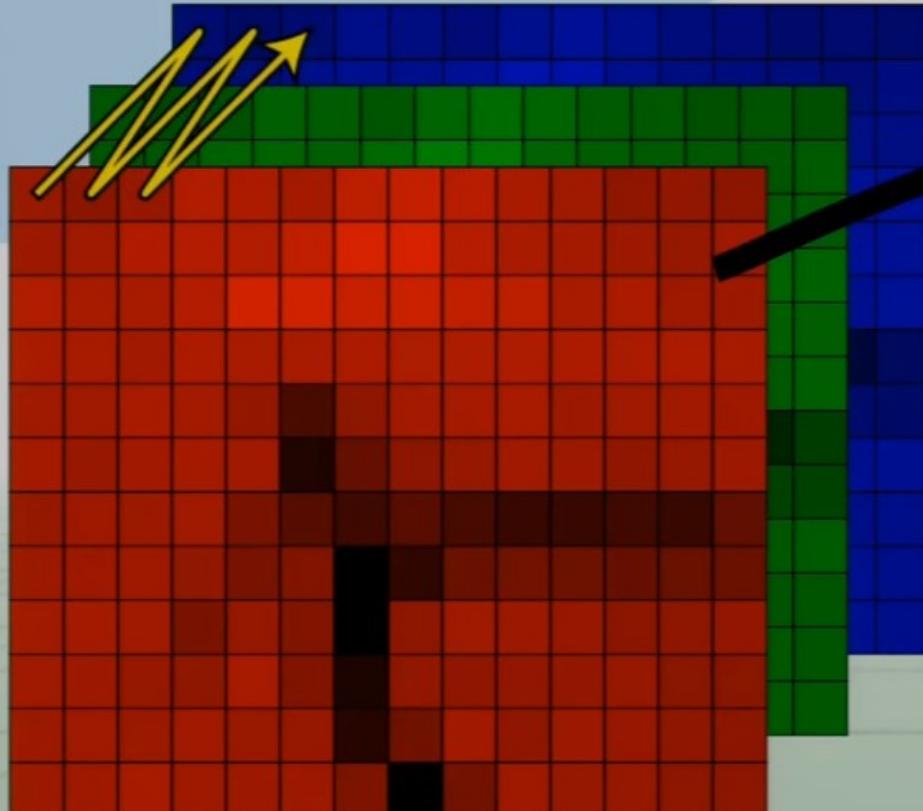
RGB information in separate “channels”

Remember: we can match “real” colors using a mix of primaries.

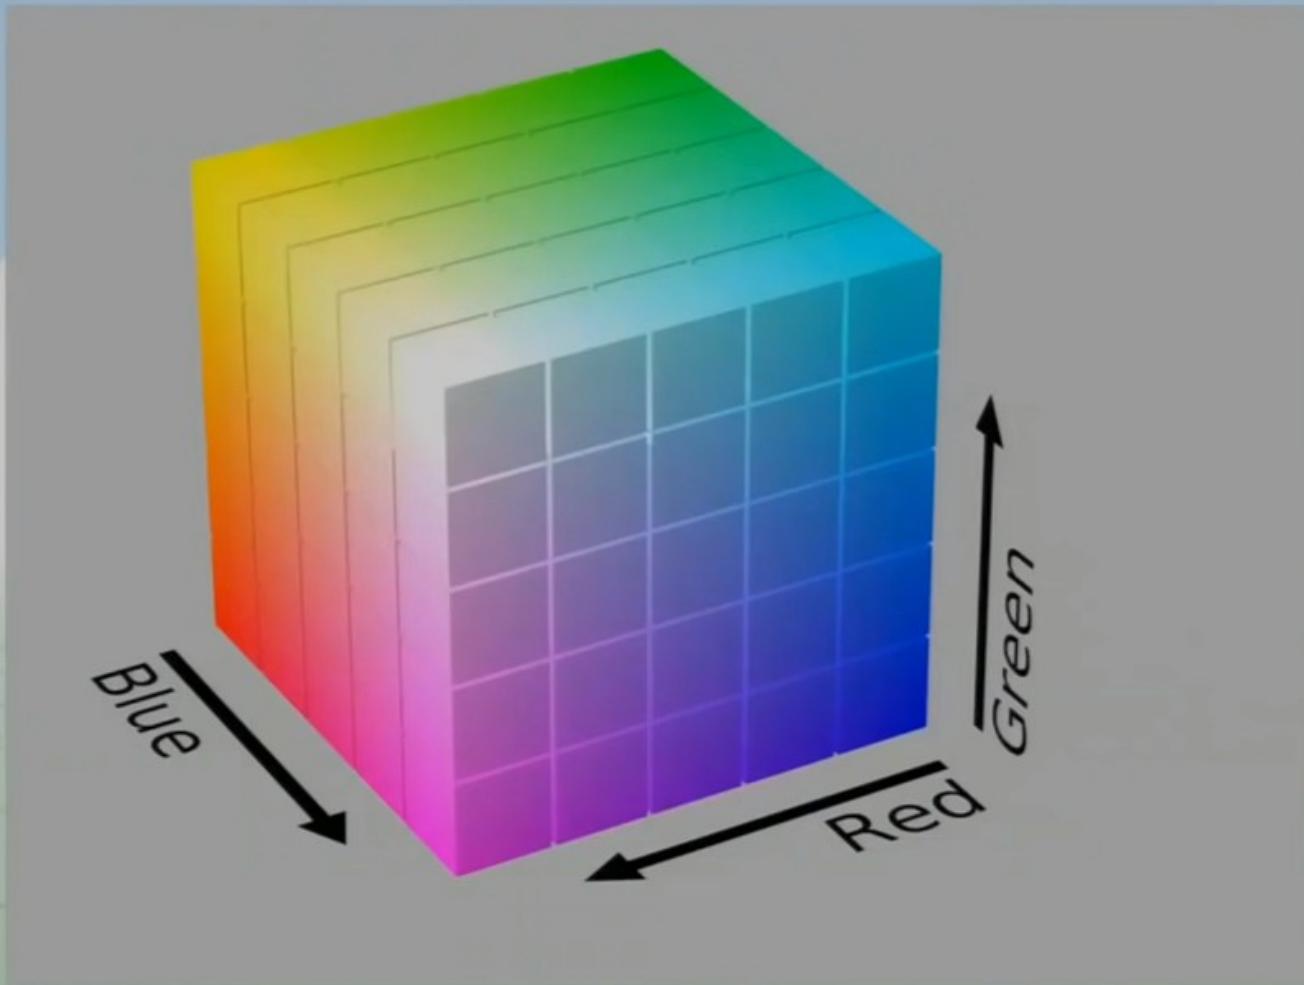
Each channel encodes one primary. Adding the light produced from each primary mimics the original color.



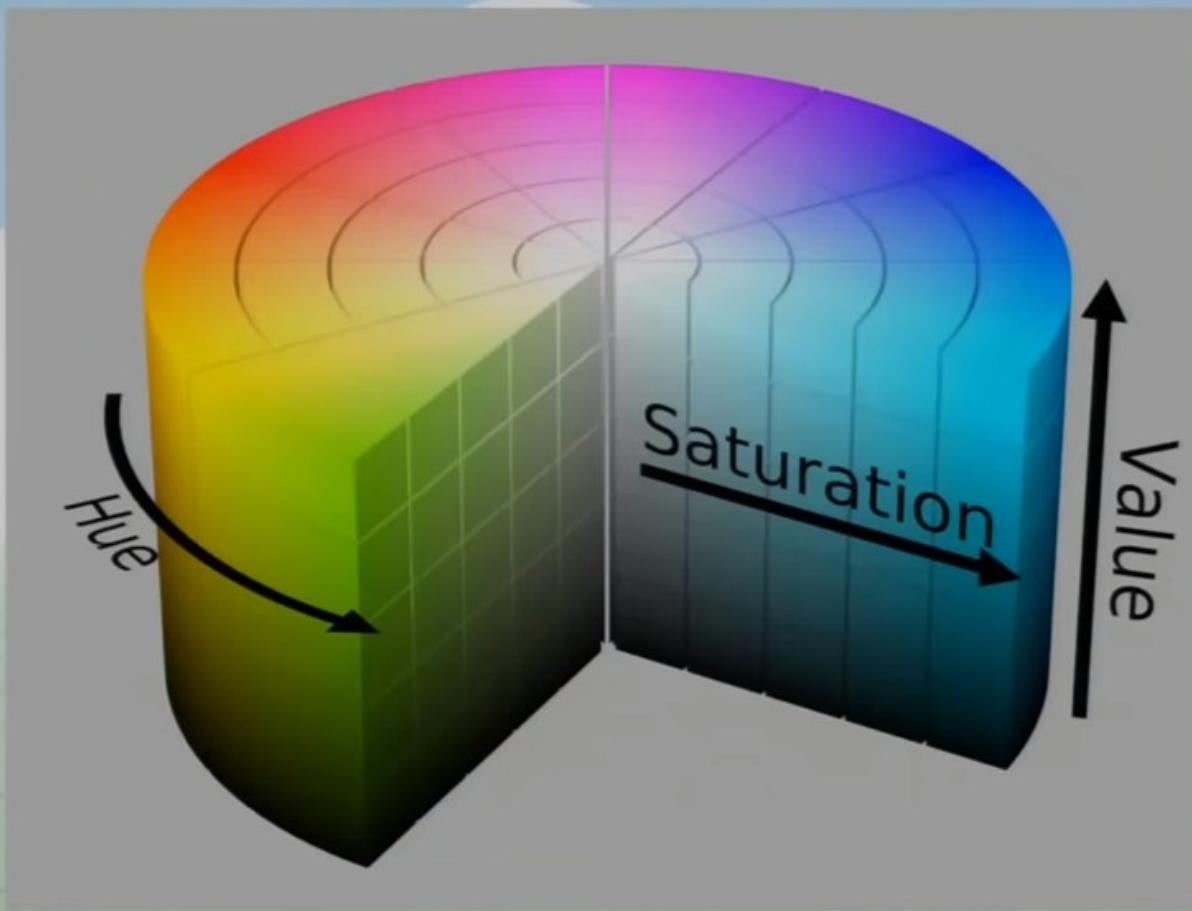
HWC: channels interleaved



RGB is a cube...

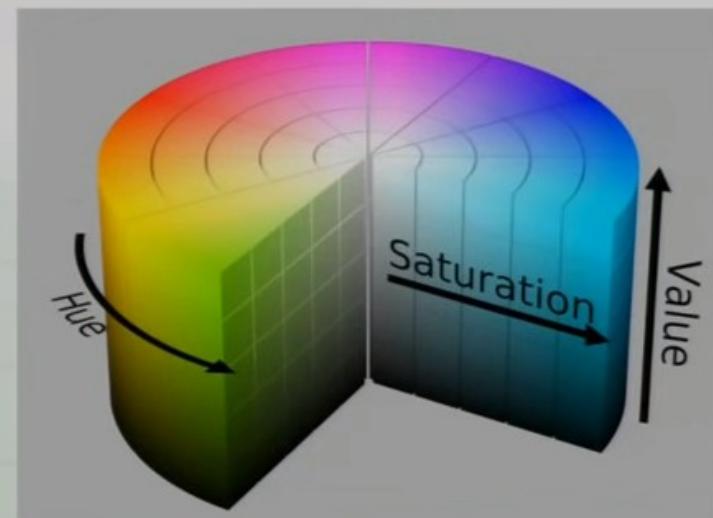


Hue, Saturation, Value: cylinder!



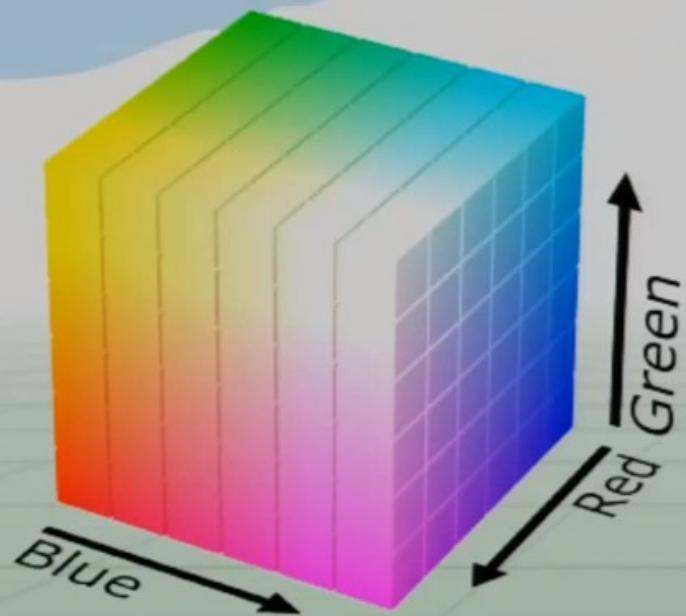
Hue, Saturation, Value

- Different model based on perception of light
- Hue: what color
- Saturation: how much color
- Value: how bright
- Allows easy image transforms
 - Shift the hue
 - Increase saturation

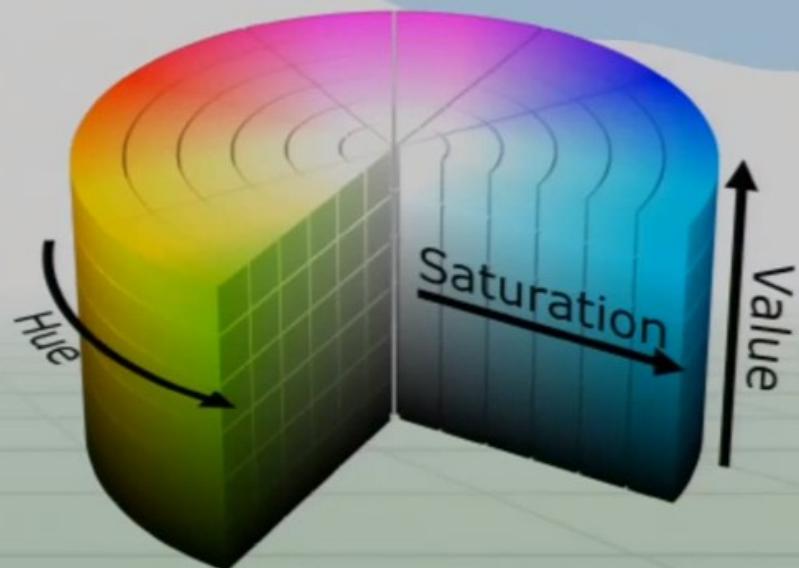


Other colorspaces are fun!

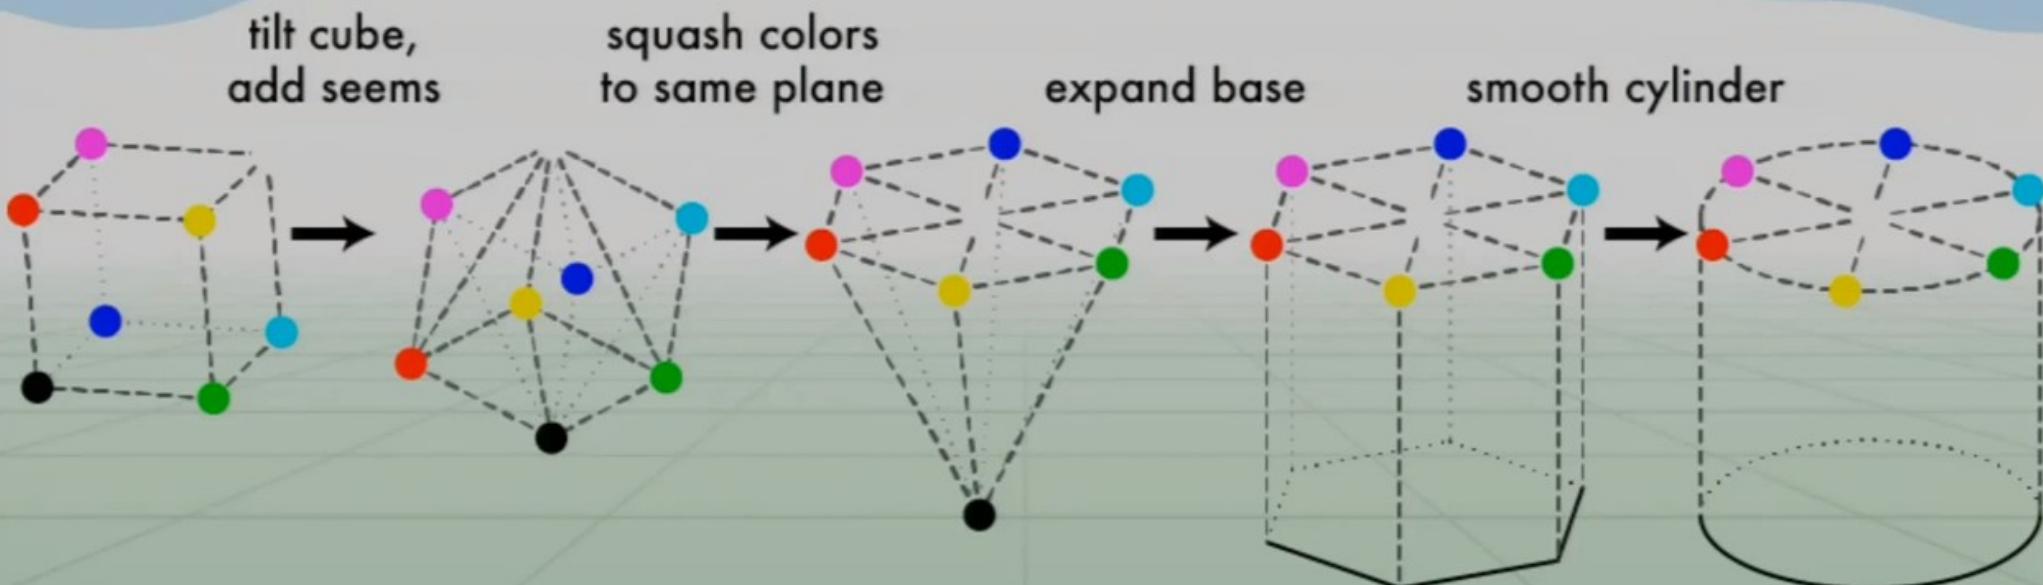
RGB



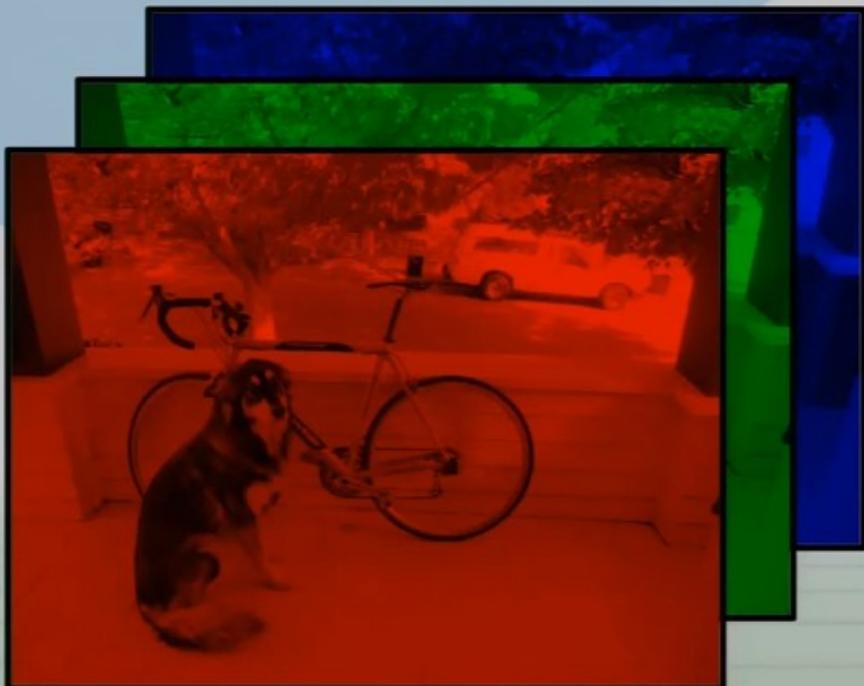
HSV



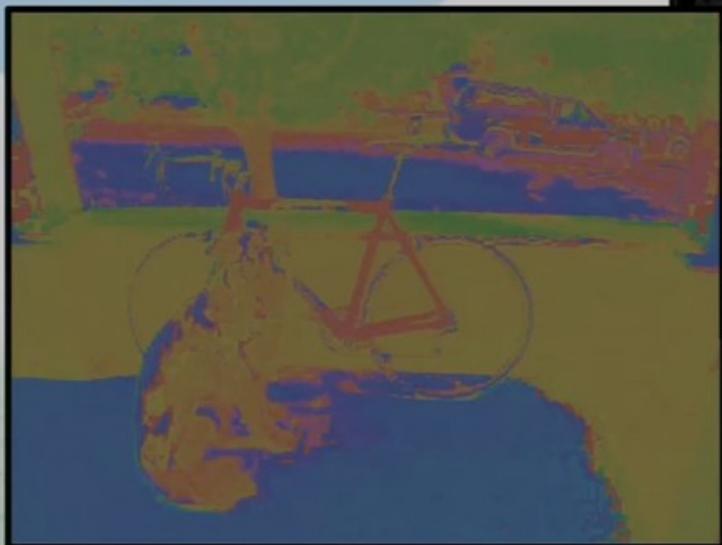
Geometric HSV to RGB:



Still 3d tensor, different info



Hue



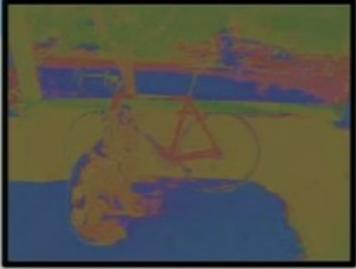
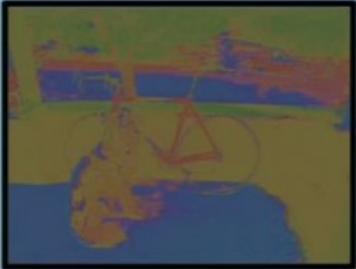
Saturation



Value



More saturation = intense colors



2x



Shift hue = shift colors



- .2

